



Embedded Operating Systems

Che-Wei Chang

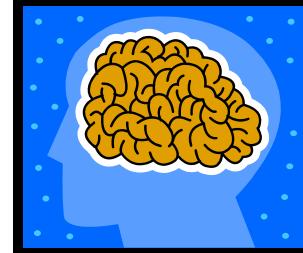
chewei@mail.cgu.edu.tw

Department of Computer Science and Information
Engineering, Chang Gung University

Course Roadmap

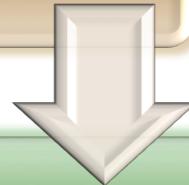
Basic Concepts

- Embedded System Design Concepts
- Embedded System Developing Tools and Operating Systems
- Embedded Linux and Android Environment



Core Technology

- Real-Time System Design and Scheduling Algorithms
- System Synchronization Protocols



Real Implementation

- System Initialization and Memory Management
- Power Management Techniques and System Routine
- Embedded Linux Labs and Exercises on Android





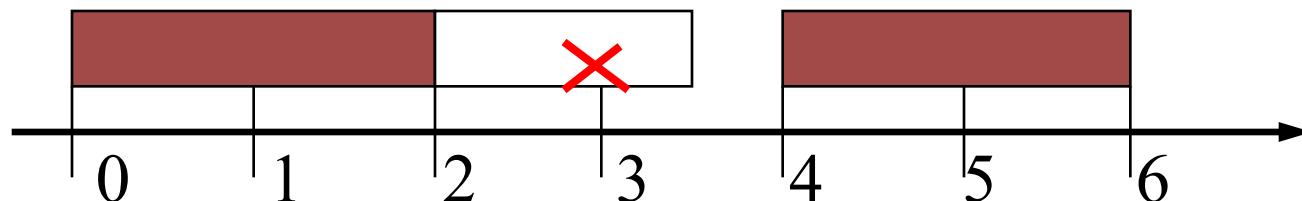
Real-Time Scheduling

Motivation

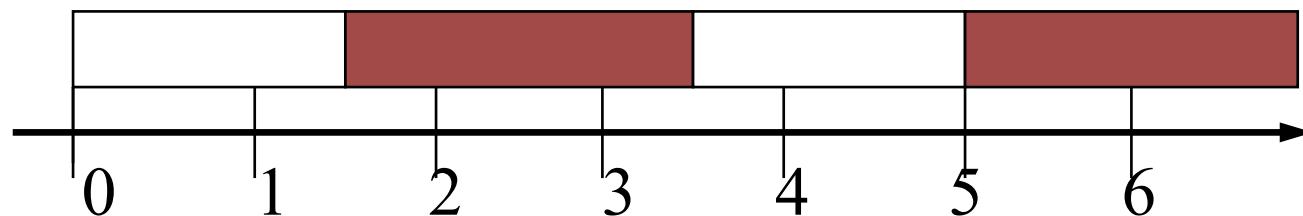
- ▶ Studying: 2 days per 4 days 

Playing Basketball: 1.5 days per 3 days 

- ▶ Case 1: Studying is always more important



- ▶ Case 2: Doing whatever is more urgent



Questions

- ▶ Can we find an **optimal** scheduler that always produces a feasible schedule whenever it is possible to do so?
 - What does **optimality** means?
- ▶ Can we find a quick schedulability test for a set of processes?
 - Is it simple and accurate?
- ▶ How do we model scheduling overheads, such as the cost of context switching?

Tentative Assumptions

- ▶ Processes are independent
- ▶ Processes are all periodic
- ▶ The deadline of a request is its next request time
- ▶ A scheduler consists of a priority assignment policy and a priority-driven scheduling mechanism

Reference: C.L. Liu and James. W. Layland, “Scheduling Algorithms for Multiprogramming in a Hard Real-Time Environment,” JACM, Vol. 20, No.1, January 1973, pp. 46-61

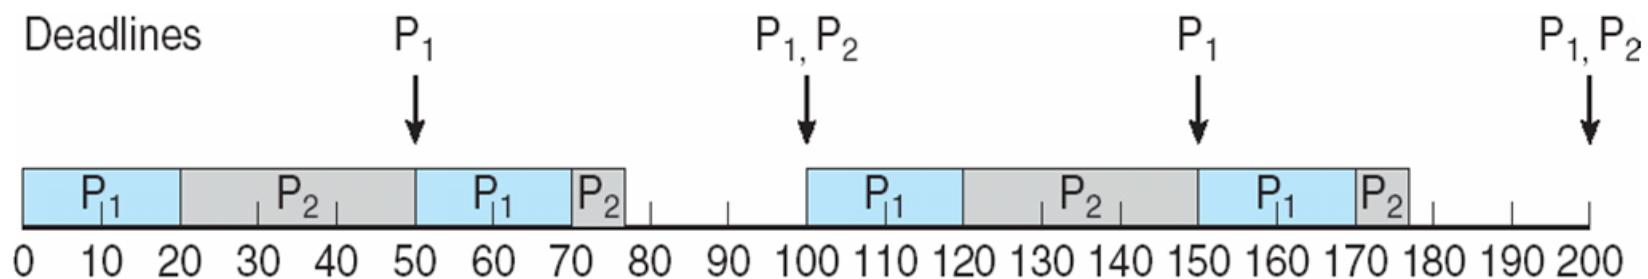
Definitions

- ▶ The **response time** of a request for a process is the time span between the request and the end of the response to that request
- ▶ A **critical instant** of a process is an instant at which a request of that process has the longest response time
- ▶ A **critical interval** for a process is the time interval between the start of a critical instant and the deadline of the corresponding request of the process
 - ➔ A critical instant for any process occurs whenever the process is requested simultaneously with requests for all higher priority processes

An observation: If a process can complete its execution within its critical interval, it is schedulable at all time!

A Static Scheduling Algorithm— Rate Monotonic Scheduling

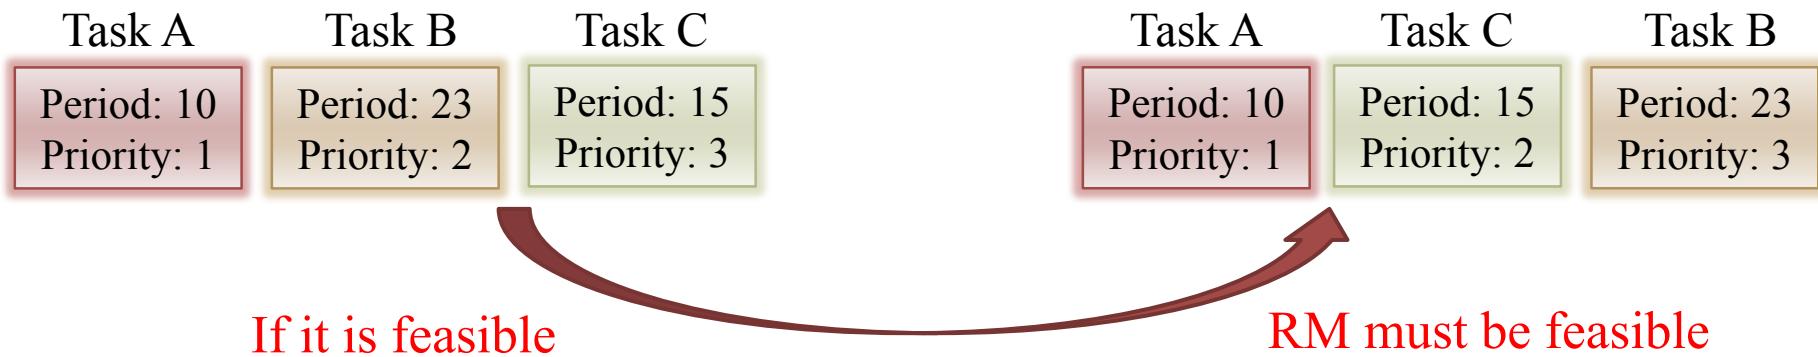
- ▶ A static priority is assigned to each task based on the inverse of its period
 - A task with shorter period → higher priority
 - A task with longer period → lower priority
 - For example:
 - P_1 has its **period 50** and execution time 20
 - P_2 has its **period 100** and execution time 37
 - P_1 is assigned a higher priority than P_2



Property of Rate Monotonic Scheduling

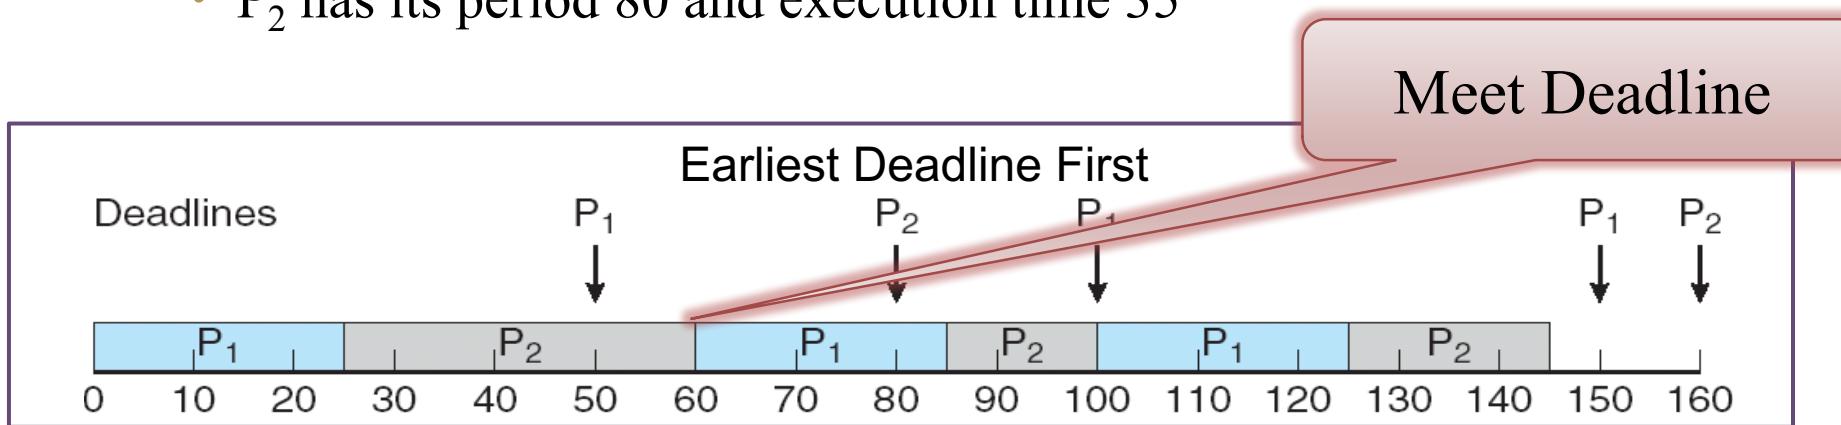
- ▶ The **rate monotonic** (RM) priority assignment assigns processes priorities according to their request rates
 - If a feasible fixed priority assignment exists for some process set, then the rate monotonic priority assignment is feasible for that process set
-  **The optimal fixed priority assignment**

Proof. Exchange the priorities of two tasks if their priorities are out of RMS order.



A Dynamic Scheduling Algorithm—Earliest Deadline First Scheduling

- ▶ Dynamic priorities are assigned according to deadlines
 - The earlier the deadline, the higher the priority
 - The later the deadline, the lower the priority
 - For example:
 - P_1 has its period 50 and execution time 25
 - P_2 has its period 80 and execution time 35



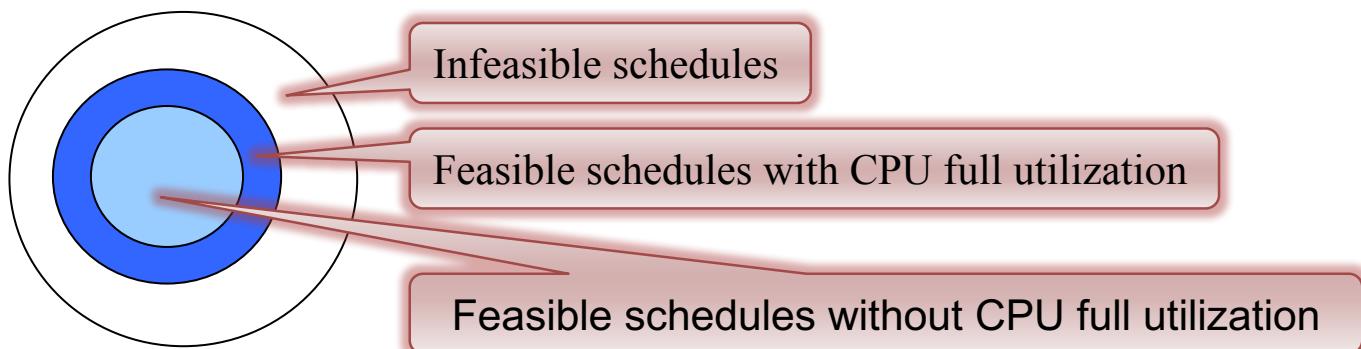
Real-Time Analysis

- ▶ For a task τ_i with the period P_i and the execution time C_i , the utilization U_i of τ_i is defined as $U_i = \frac{C_i}{P_i}$
- ▶ For a real-time task set T the total utilization of the task set is $\sum_{\tau_i \in T} U_i$
- ▶ If $\sum_{\tau_i \in T} U_i \leq 69\%$, Rate Monotonic Scheduling can schedule all tasks in T to meet all deadlines
 - More precisely, for n tasks, the i -th task can meet deadline if
$$\sum_{j=1}^i U_j \leq i(2^{1/i} - 1)$$
- ▶ If and only if $\sum_{\tau_i \in T} U_i \leq 100\%$, Earliest Deadline First Scheduling can schedule all tasks in T to meet all deadlines

Reference: C.L. Liu and James. W. Layland, "Scheduling Algorithms for Multiprogramming in a Hard Real-Time Environment," JACM, Vol. 20, No.1, January 1973, pp. 46-61

CPU Utilization

- ▶ For a given priority assignment, a process set **fully utilizes** the processor if the priority assignment is feasible for the set and if any increase in the run time of any processes in the set will make the priority assignment **infeasible**
 - EDF: 100% → fully utilize, <100% → not fully utilize
 - RM:



RM and EDF (1 / 2)

- ▶ The achievable utilization factor of the EDF algorithm is 100%. The EDF algorithm is an optimal dynamic priority scheduling policy in the sense that a process set is schedulable if its CPU utilization is no larger than 100%.
- ▶ The achievable utilization factor of the RMS algorithm is about $\ln 2$ (~69%). The RMS algorithm is an optimal fixed priority scheduling policy in the sense that if a process set is schedulable by some fixed priority scheduling algorithm, then it is schedulable by the RMS algorithm.

RM and EDF (2/2)

- ▶ For a set of m processes with the RM fixed priority order, the i-th process is schedulable if

$$\sum_{j=1}^i \frac{c_j}{p_j} \leq i(2^{1/i} - 1)$$

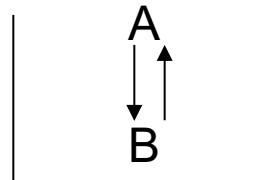
- ▶ For a set of m processes with the EDF scheduling, all process will miss deadlines when the total utilization is more than 100%

Scheduling Overheads

► Context Switching

- Needed either when a process is preempted by another process, or when a process completes its execution
- Stack Discipline

If process A preempts process B, process A must complete before process B can resume

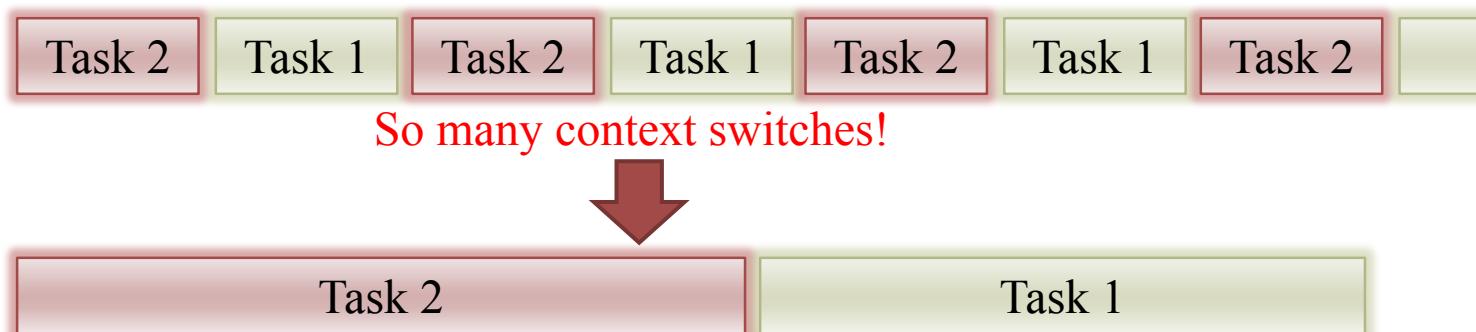


If it is obeyed, charge the cost of preemption (context switching cost) once to the preempting process!



Least Slack Time Algorithm

- ▶ The least slack time algorithm (LST), which assigns processes priorities inversely proportional to their slack times is also optimal if context switching cost can be ignored
 - The slack time of a process is $d(t) - t - c(t)$
 - t : current time
 - $d(t)$: deadline
 - $c(t)$: remaining execution time
 - An example
 - The time $t = 0$
 - Two task have the same deadline 20
 - Task 1 has $c(t) = 7$, and task 2 has $c(t) = 8$





Process Synchronization

Basic Concept

- ▶ Processes might share non-preemptible resources or have precedence constraints
- ▶ Papers for discussion:
 - L. Sha, R. Rajkumar, J.P. Lehoczky, “Priority Inheritance Protocols: An Approach to Real-Time Synchronization,” IEEE Transactions on Computers, 1990.
 - A.K. Mok, “The Design of Real-Time Programming Systems Based on Process Models,” IEEE Real-Time Systems Symposium, Dec 1994.

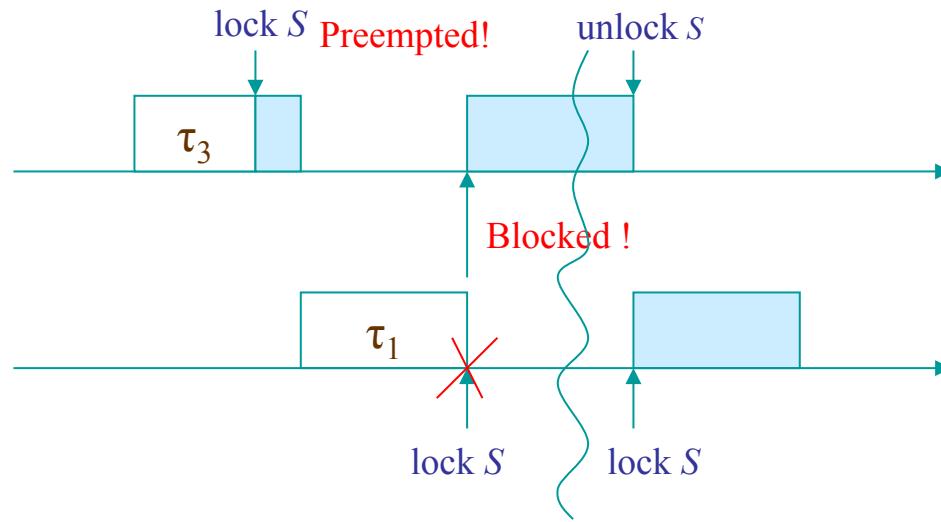
Process Synchronization

► Motivation

- Can we find an efficient way to analyze the schedulability of a process set (systematically)
- What kinds of restrictions on the use of communication primitives are needed so as to efficiently solve the restricted scheduling problem
- How can we control the priority inversion problem
- The lengths of critical sections might be quite different

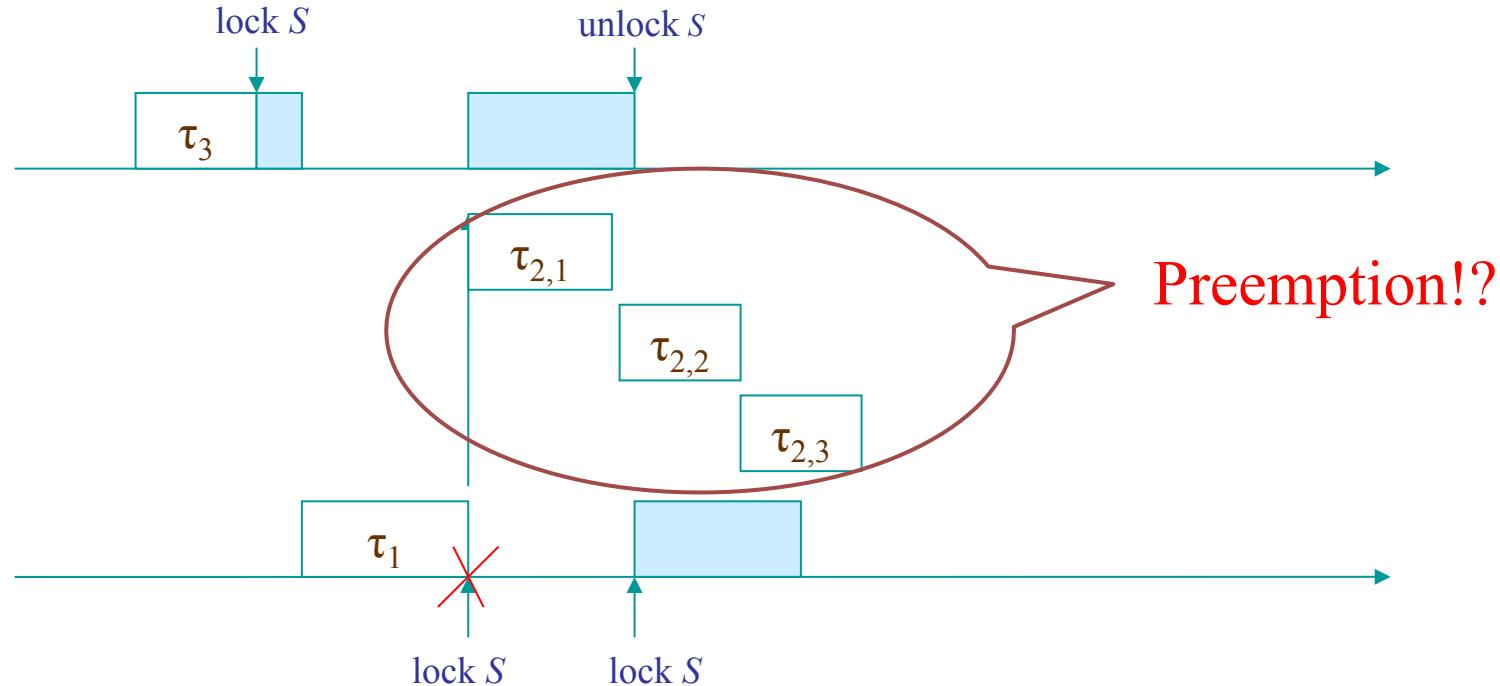
Blocking and Preemption

- ▶ Blocking: a higher-priority process is forced to wait for the execution of a lower-priority process
- ▶ Preemption: a low-priority process is forced to wait for the execution of a high-priority process



Priority Inversion

- When there are a lot of tasks having priority between that of τ_1 and τ_3 , there are a lot of priority inversions

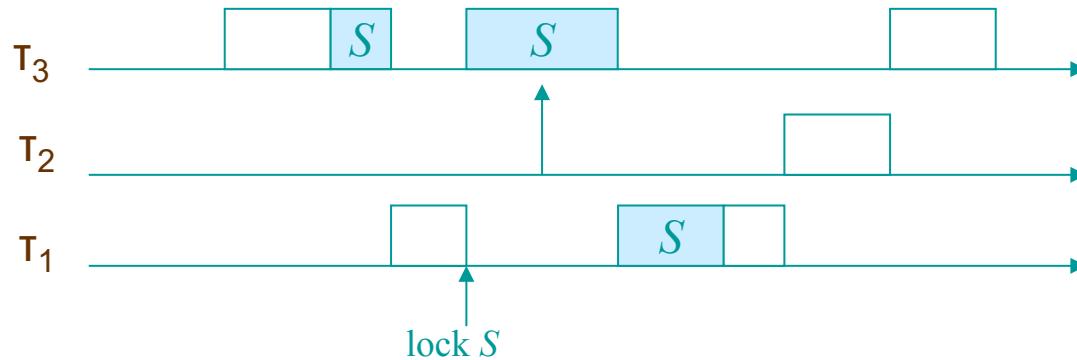


Priority Inheritance Protocol (PIP)

- ▶ Priority-Driven Scheduling
 - The process which has the highest priority among the ready processes is assigned the processor
- ▶ Synchronization
 - Process τ_i must obtain the lock on the semaphore guarding a critical section before τ_i enters the critical section
 - If τ_i obtains the required lock, τ_i enters the corresponding critical section; otherwise, τ_i is blocked and said to be blocked by the process holds the lock on the corresponding semaphore
 - Once τ_i exits a critical section, τ_i unlocks the corresponding semaphore and makes its blocked processes ready
- ▶ Priority Inheritance
 - If a process τ_i blocks higher priority processes, τ_i inherits the highest priority of the process blocked by τ_i
 - Priority inheritance is transitive

Properties of PIP

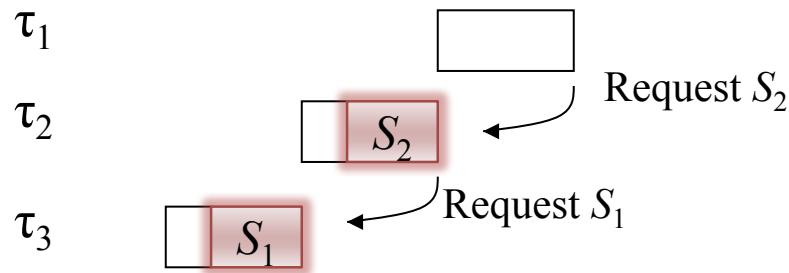
- ▶ No priority inversion



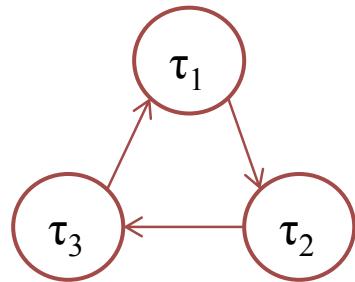
- ▶ A semaphore S can be used to cause inheritance blocking to task J only if S is accessed by a task which has a priority lower than that of J and might be accessed by a task which has a priority equal to or higher than that of J .

Concerns of PIP

- ▶ A chain of blocking is possible



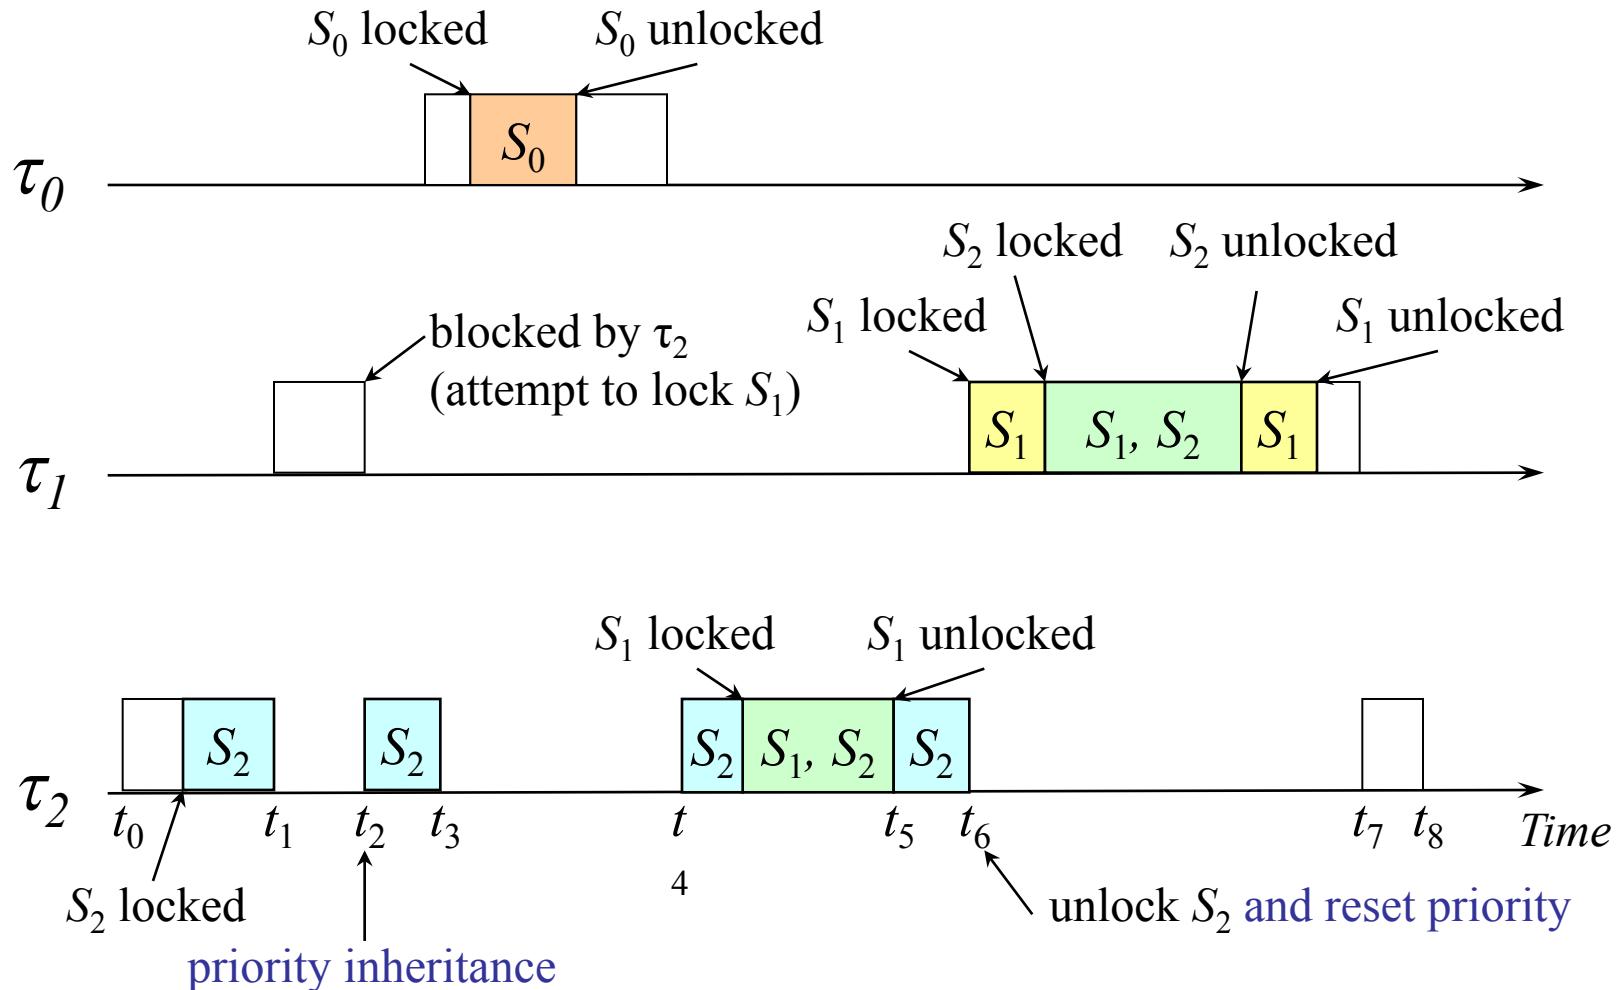
- ▶ A deadlock can be formed



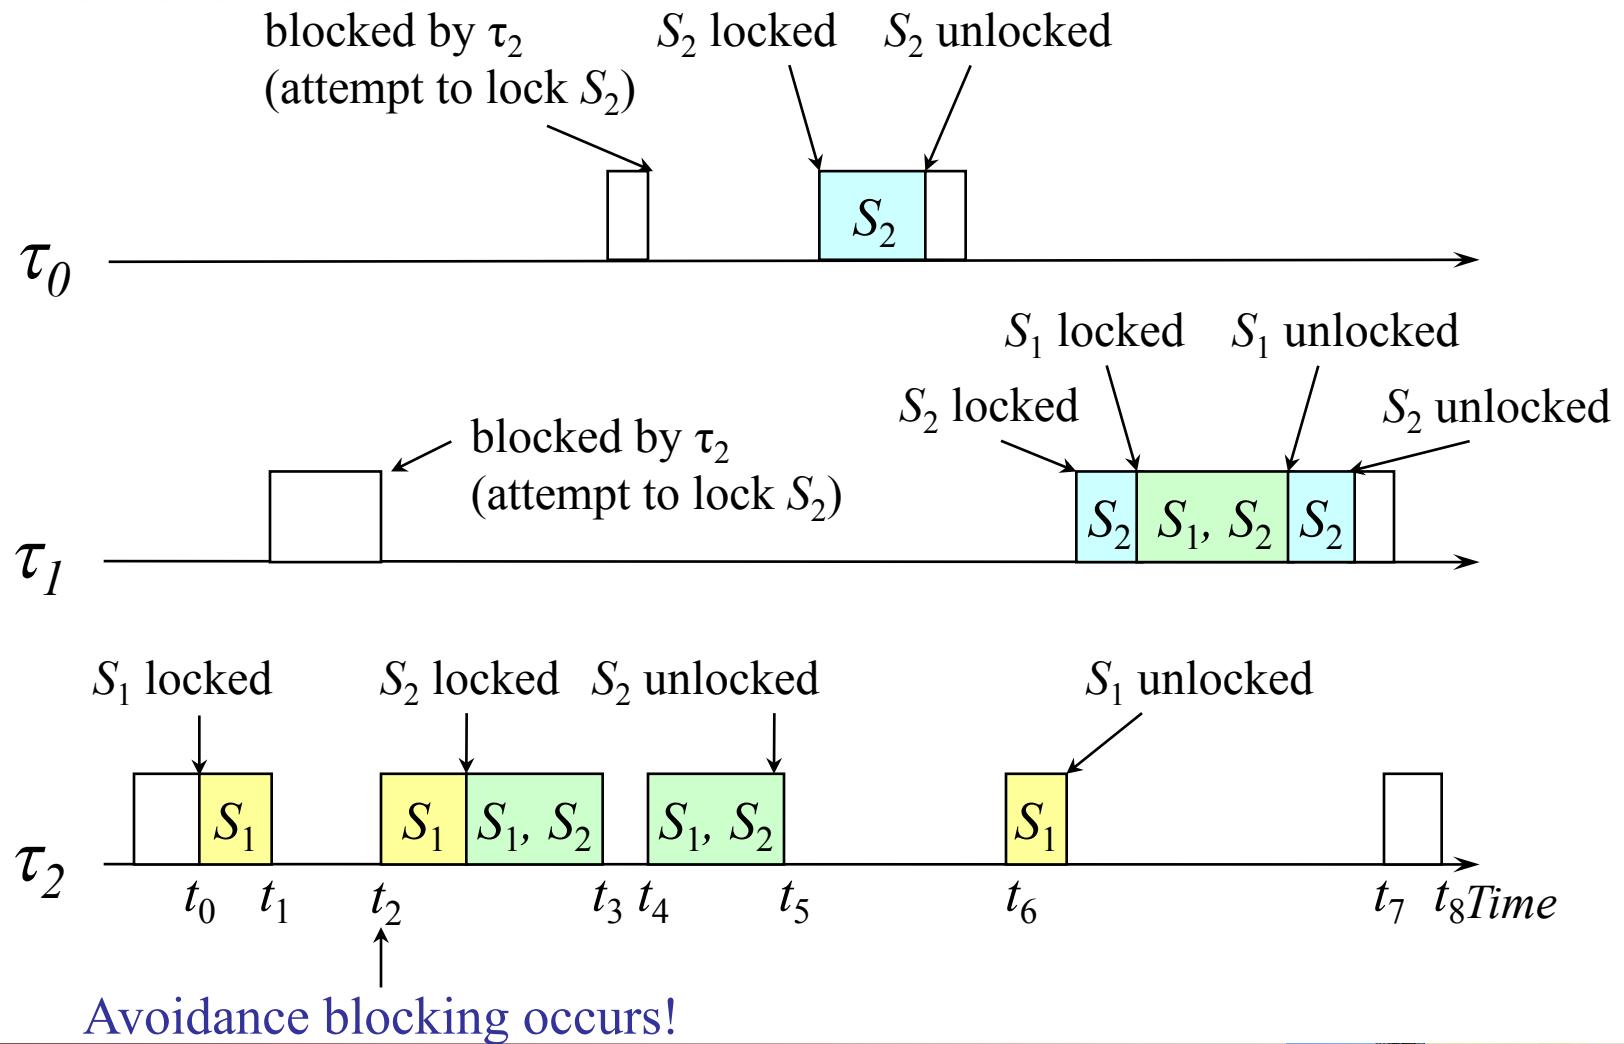
Priority Ceiling Protocol (PCP)

- ▶ The priority ceiling of a semaphore is the priority of the highest priority task that may lock the semaphore
- ▶ The Basic Priority Inheritance Protocol + Priority Ceiling
- ▶ A task J may successfully lock a semaphore S if S is available, and the priority of J is higher than the highest priority ceiling of all semaphores currently locked by tasks other than J
- ▶ Priority inheritance is transitive

Example: Deadlock Avoidance



Example: Chain Blocking Avoidance



Properties of PCP

- ▶ The priority ceiling protocol prevents transitive blockings
- ▶ The priority ceiling protocol prevents deadlock
- ▶ No job can be blocked for more than one critical section of any lower priority job
- ▶ A set of n periodic tasks under the **priority ceiling protocol** can be scheduled by the **rate monotonic algorithm** if the following conditions are satisfied:

$$\forall i, \quad 1 \leq i \leq n, \quad \sum_{j=1}^{i-1} \frac{c_j}{p_j} + \frac{c_i + B_i}{p_i} \leq i(2^{1/i} - 1),$$

where B_i is the worst-case blocking time for τ_i , and each task will be blocked on once in a period



Rate Monotonic Analysis

Periodic Requirements (1 / 2)

Task τ_1 : $C_1=20$, $P_1=100$, $U_1=0.2$

Task τ_2 : $C_2=40$, $P_2=150$, $U_2=0.267$

Task τ_3 : $C_3=100$, $P_3=350$, $U_3=0.286$

- ▶ Total utilization: $75.3\% \leq 3 \left(2^{\frac{1}{3}} - 1 \right) = 77.9\%$
- ▶ 24.7% of the CPU is usable for lower-priority background computation

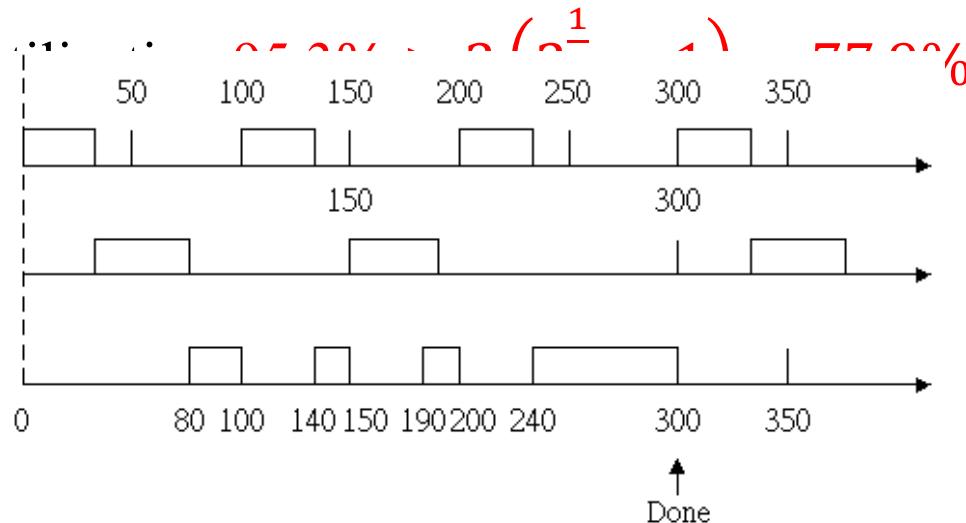
Periodic Requirements (2/2)

Task τ_1 : $C_1=40$, $P_1=100$, $U_1=0.4$

Task τ_2 : $C_2=40$, $P_2=150$, $U_2=0.267$

Task τ_3 : $C_3=100$, $P_3=350$, $U_3=0.286$

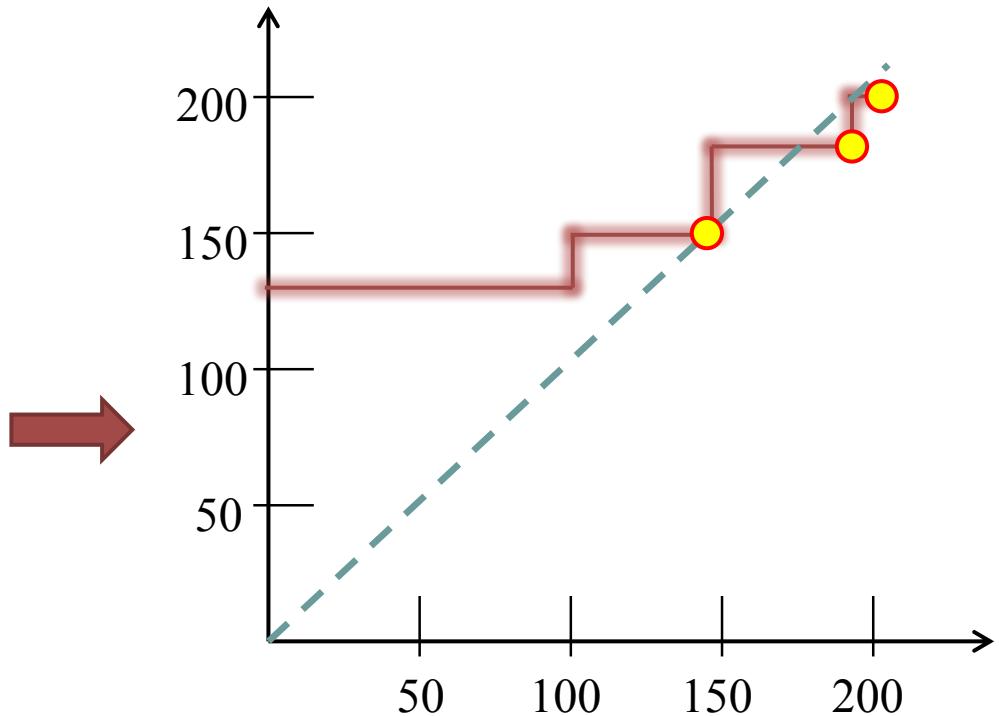
- ▶ The utilization of the first two tasks: $66.7\% \leq 2 \left(2^{\frac{1}{2}} - 1 \right) = 82.8\%$
- ▶ The total



Rate Monotonic Analysis (RMA)

► A RMA Example:

- $\tau_1(20,100)$, $\tau_2(30,150)$, $\tau_3(80, 210)$, $\tau_4(100,400)$
- τ_1
 - $c_1 \leq 100$
- τ_2
 - $c_1 + c_2 \leq 100$ or
 - $2c_1 + c_2 \leq 150$
- τ_3
 - $c_1 + c_2 + c_3 \leq 100$ or
 - $2c_1 + c_2 + c_3 \leq 150$ or
 - $2c_1 + 2c_2 + c_3 \leq 200$ or
 - $3c_1 + 2c_2 + c_3 \leq 210$
- τ_4
 - $c_1 + c_2 + c_3 + c_4 \leq 100$ or
 - $2c_1 + c_2 + c_3 + c_4 \leq 150$ or
 - ...



RMA with Blocking Consideration (1 / 2)

- ▶ A RMA Example with blocking time:
 - $\tau_1(20,100), \tau_2(30,150), \tau_3(80, 210), \tau_4(100,400)$
 - $\tau_1: (S_1, 5)$
 - $\tau_2: (S_2, 15)$
 - $\tau_3: (S_1, 10), (S_3, 5)$
 - $\tau_4: (S_2, 5), (S_3, 20)$
- ▶ What is the priority ceiling of each semaphore?
 - $S_1: \tau_1, S_2: \tau_2, S_3: \tau_3$
- ▶ When PCP is adopted (block once), what is the blocking time of each task?
 - $\tau_1: 10, \tau_2: 10, \tau_3: 20, \tau_4: 0$

RMA with Blocking Consideration (2/2)

- ▶ A RMA Example with blocking time:
 - For each task, we have to consider the execution time, period, and blocking time
 - $\tau_1(20,100,10)$, $\tau_2(30,150,10)$, $\tau_3(80, 210, 20)$, $\tau_4(100,400,0)$
 - τ_1
 - $b_1 + c_1 \leq 100$
 - τ_2
 - $b_2 + c_1 + c_2 \leq 100$ or
 - $b_2 + 2c_1 + c_2 \leq 150$
 - τ_3
 - $b_3 + c_1 + c_2 + c_3 \leq 100$ or
 - $b_3 + 2c_1 + c_2 + c_3 \leq 150$ or
 - $b_3 + 2c_1 + 2c_2 + c_3 \leq 200$ or
 - $b_3 + 3c_1 + 2c_2 + c_3 \leq 210$
 - τ_4
 - ...



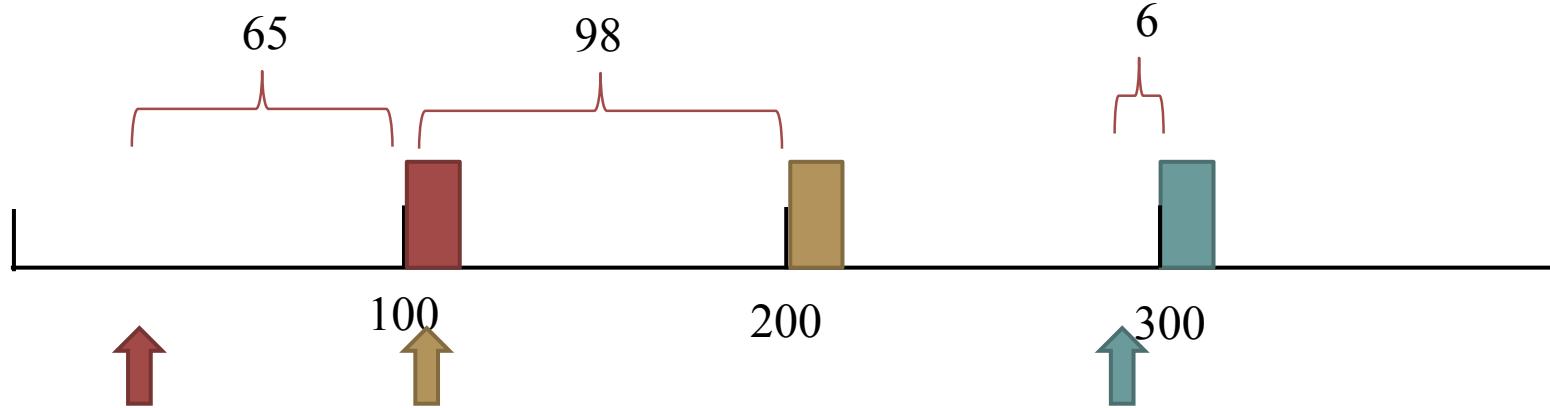
Aperiodic Servers

Observation of Aperiodic Tasks

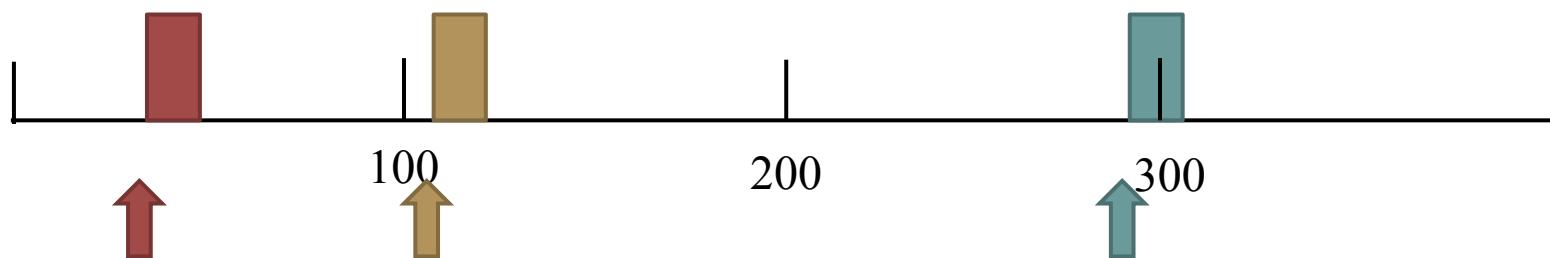
- ▶ Aperiodic tasks run at irregular intervals
- ▶ Aperiodic deadlines
 - Hard deadline: minimum inter-arrival time
 - Soft deadline: best average response time
- ▶ Services such as
 - User requests
 - Device interrupts
 - ...

Scheduling Aperiodic Tasks

- ▶ Polling Server ~ Average Response Time = 50 units

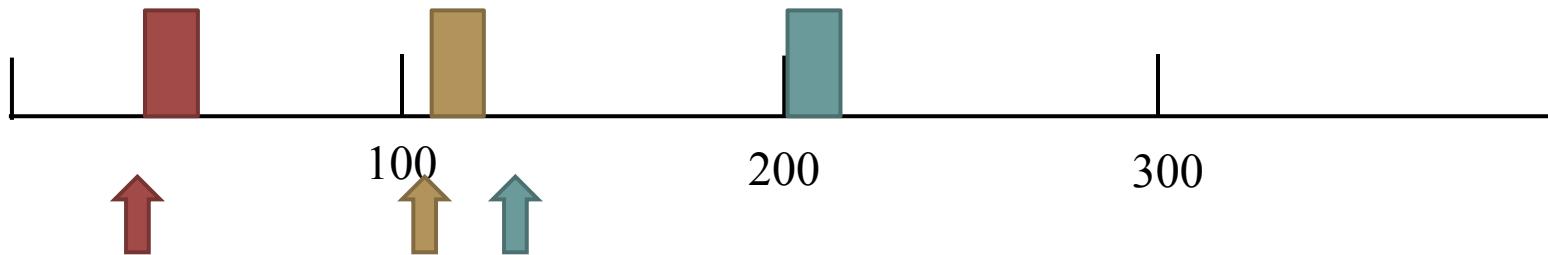


- ▶ Interrupt Server ~ Average Response Time = 1 unit



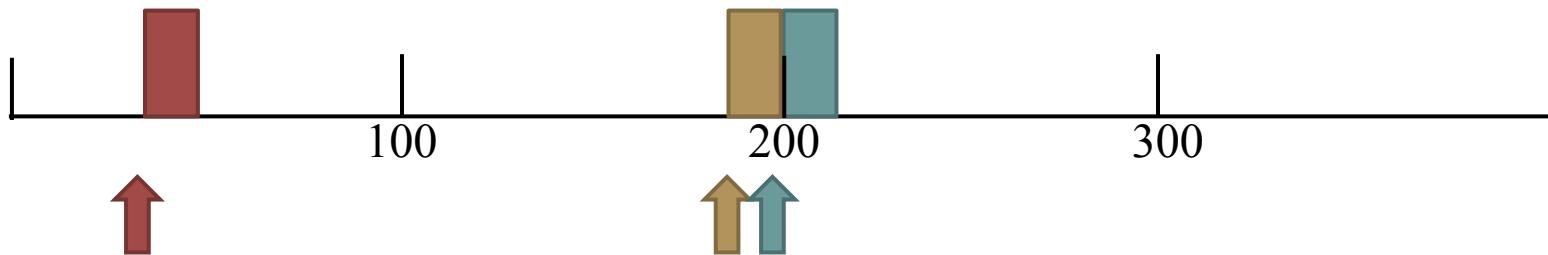
Deferrable Server

- ▶ Polling Server: the average response time is long
- ▶ Interrupt Server: the computing time of aperiodic tasks is difficult to limited
- ▶ Deferrable Server
 - In each period, a deferrable server has a execution budget
 - When execution budget is used up, server execution drops to a lower (background) priority

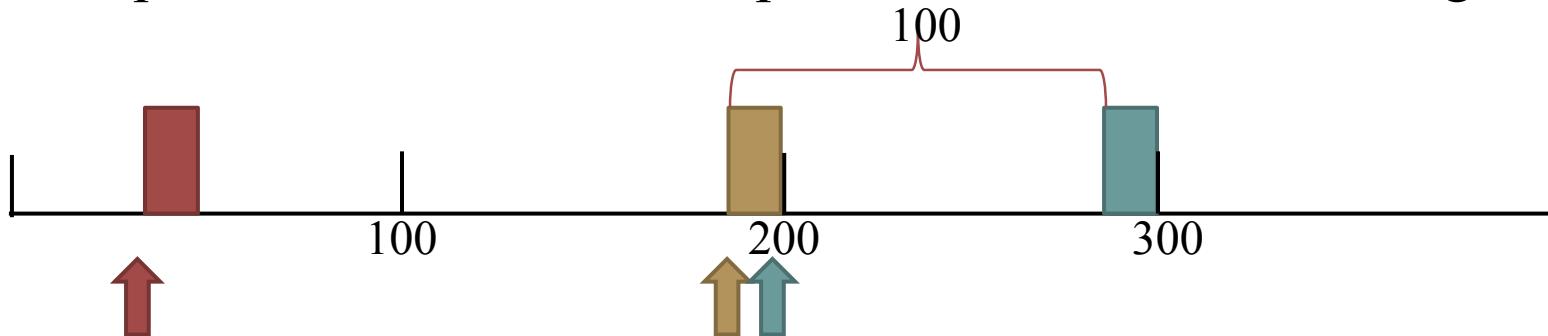


Sporadic Server

- ▶ Deferrable Server might consume two times of the execution budget in short time



- ▶ Sporadic Server
 - Replenishment occurs one “period” after the start of usage

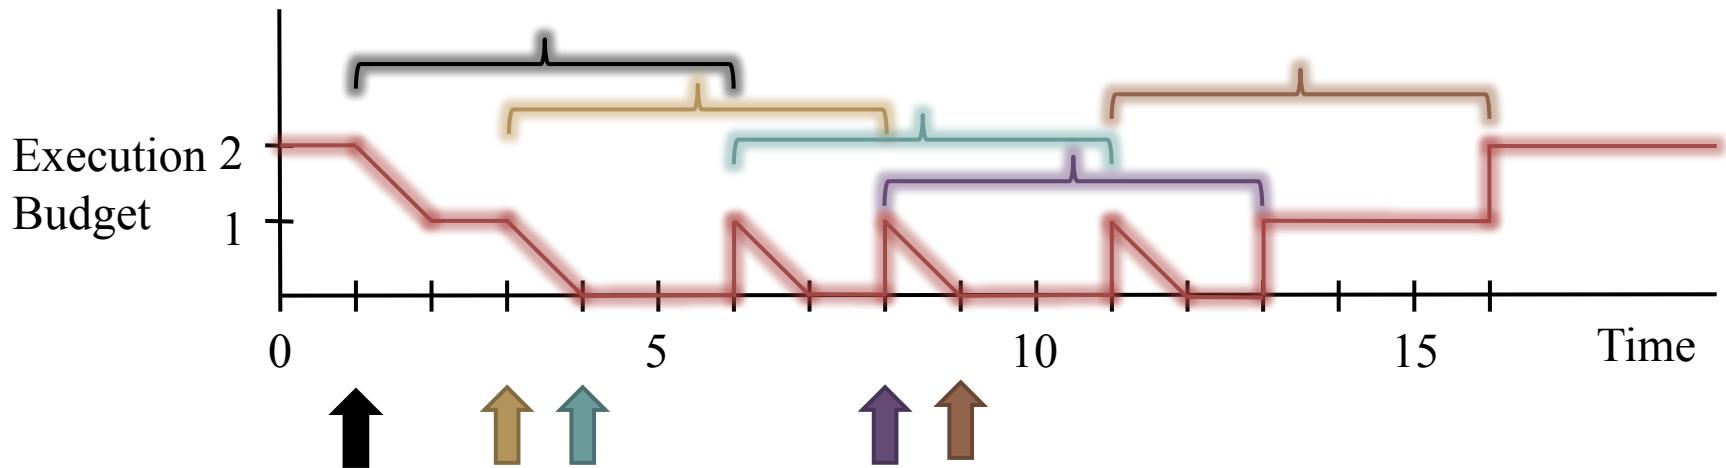


Properties of Sporadic Server

- ▶ A sporadic server differs from a deferrable server in its replenishment policy:
 - A 100 ms deferrable server replenishes its execution budget every 100 ms, no matter when the execution budget is used
 - The affect of a sporadic server on lower priority tasks is no worse than a periodic task with the same period and execution time

An Example of Sporadic Server

- ▶ A sporadic server has a replenishment period 5 and an execution budget 2
- ▶ Each event consumes the execution 1
- ▶ Events arrive at 1, 3, 4, 8, 9



Properties of Sporadic Server

- ▶ For a sporadic server has a replenishment period X and an execution budget Y
 - Given a set of sporadic tasks, If
 - Each of the aperiodic tasks has its minimum inter-arrival time no less than X
 - The total execution of the task set is no more than Y
 - All sporadic tasks can meet the deadline constraints
- ▶ When a system consists of periodic tasks and sporadic servers
 - A sporadic server with replenishment period X and an execution budget Y can be consider as a periodic task with a period X and an execution time Y
 - The system can then use analysis scheme of RM or EDF