



Embedded Operating Systems

Che-Wei Chang

chewei@mail.cgu.edu.tw

Department of Computer Science and Information
Engineering, Chang Gung University

Grading and Resources

- ▶ Midterm: 30%
- ▶ Discussion and Quiz: 20%
- ▶ Exercises: 20%
- ▶ Final Project Presentation and Report: 30%
- ▶ Course Website:
<http://www.csie.cgu.edu.tw/~chewei/teaching.html>
- ▶ Moved to
<https://icechewei.github.io/webpage/teaching.html>

Paper Reading and Topic Survey

- ▶ You can cover all contents of the paper or survey the topic by yourself
- ▶ 10 minutes for each presentation
- ▶ 5 extra minutes for Q&A
- ▶ Number of slides is from 10 to 20

Reading List

Paper	Name	Student ID
1 My VM is Lighter (and Safer) than your Container	M0629013	盧昱宏
2 NetCache: Balancing Key-Value Stores with Fast In-Network Caching	M0629010	呂毓軒
3 Resource Central: Understanding and Predicting Workloads for Improved Resource Management in Large Cloud Platforms	M0629012	陳志榮
4 Corey: An Operating System for Many Cores	M0529004	張孝荃
5 Scaling a file system to many cores using an operation log	M0529016	許鏡照
6 Dynamo: Amazon's Highly Available Key-value Store	M0529011	SABIH AHMAD KHAN
7 SVE: Distributed Video Processing at Facebook Scale	M0629009	沈凱明
8 NEVE: Nested Virtualization Extensions for ARM	D9921002	詹程凱
9 Multiprogramming a 64 kB Computer Safely and Efficiently	M0629014	王宥憲
10 NOVA-Fortis: A Fault-Tolerant Non-Volatile Main Memory File System	M0629018	章齊信



Case Study:

Real-Time Partitioned Scheduling on Multi-Core Systems with Local and Global Memories



Real-Time Partitioned Scheduling on Multi-Core Systems with Local and Global Memories

Che-Wei Chang, Jian-Jia Chen, Tei-Wei Kuo and Heiko Falk

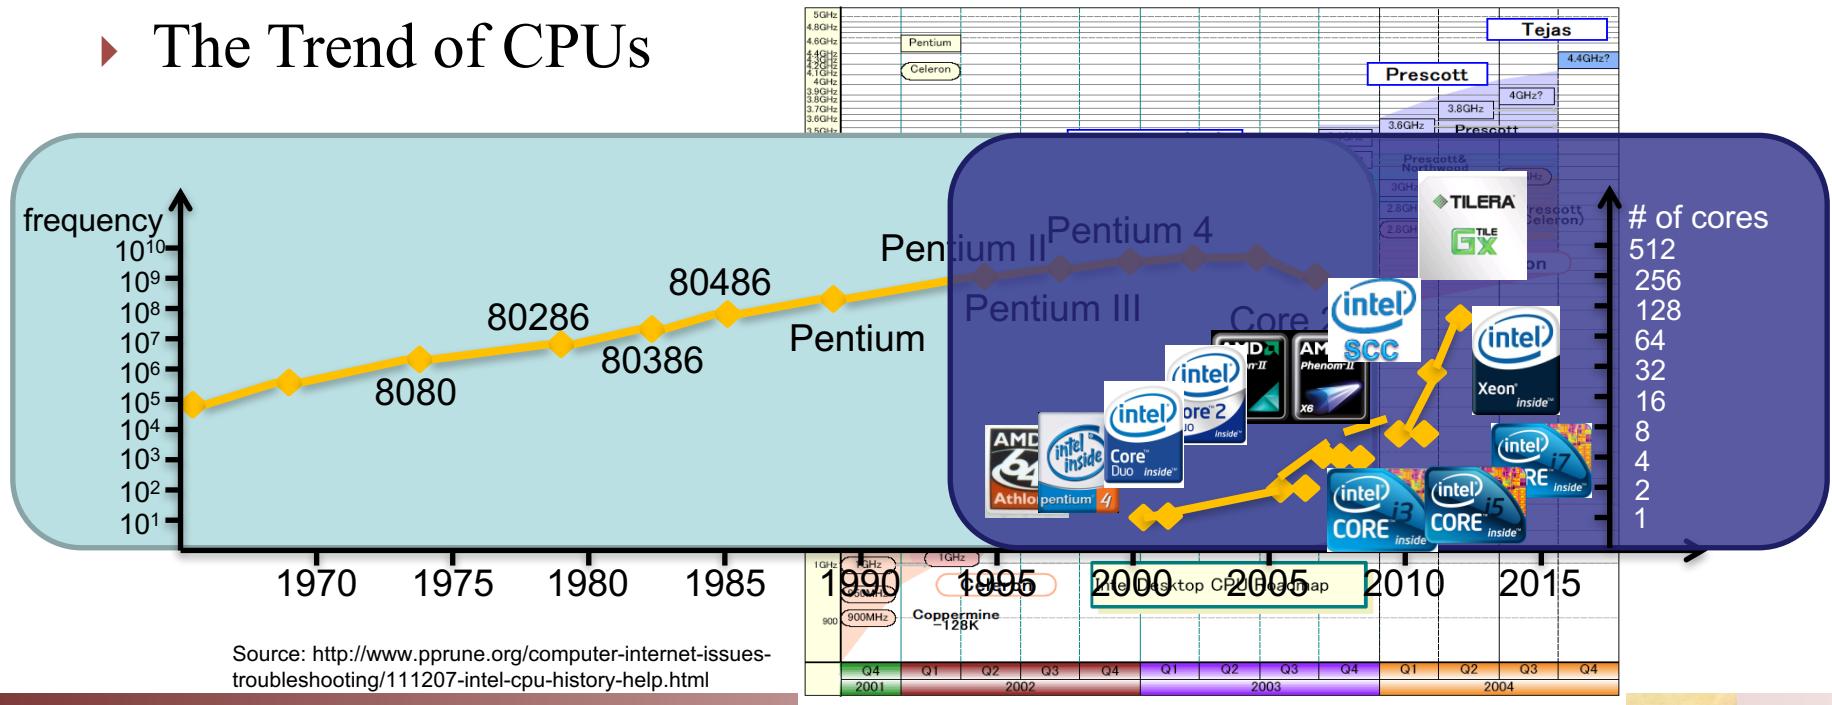
18th Asia and South Pacific Design Automation Conference, Jan. 22–25, Yokohama, Japan

Agenda

- ▶ Introduction and Motivation
- ▶ System Architecture of Multiprocessor Platforms with Local and Global Memories
- ▶ A Scheduling Algorithm with Considerations of Real-Time and Memory-Space Constraints
- ▶ Performance Evaluation with Benchmark Suites
- ▶ Conclusion

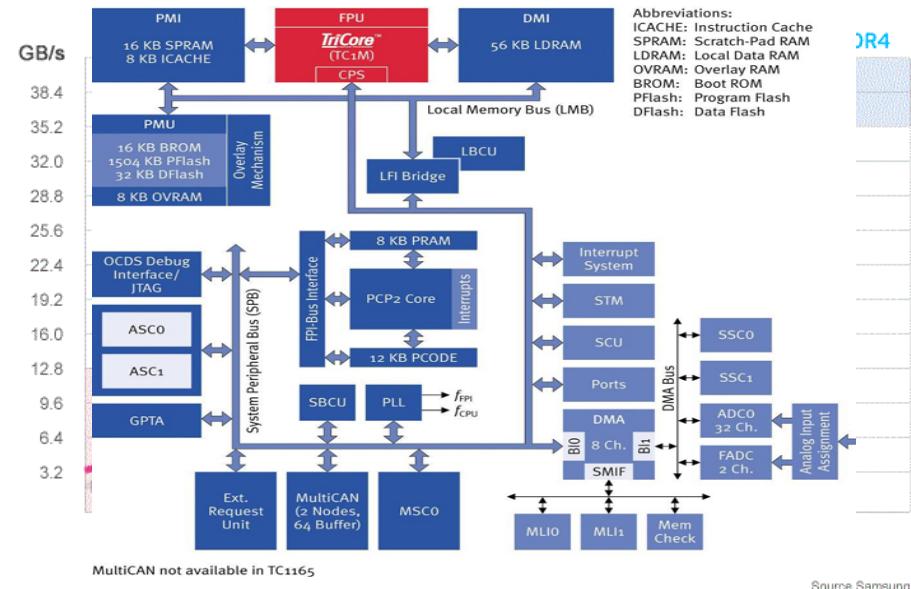
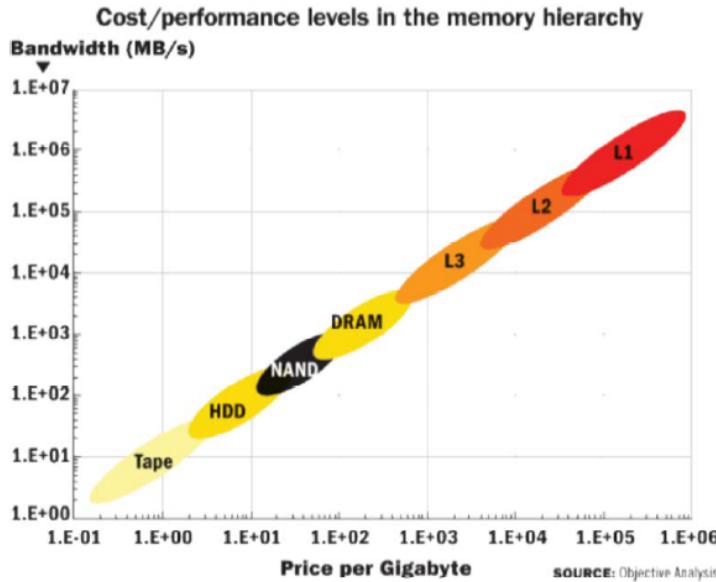
Introduction: Multi-Core Systems

- ▶ Emerging Problems of Computing System Designs
 - Limited battery power for embedded systems
 - Thermal problems of integrated chips
 - Frequency constraints of high-performance processors
- ▶ The Trend of CPUs



Motivation: Memory Wall

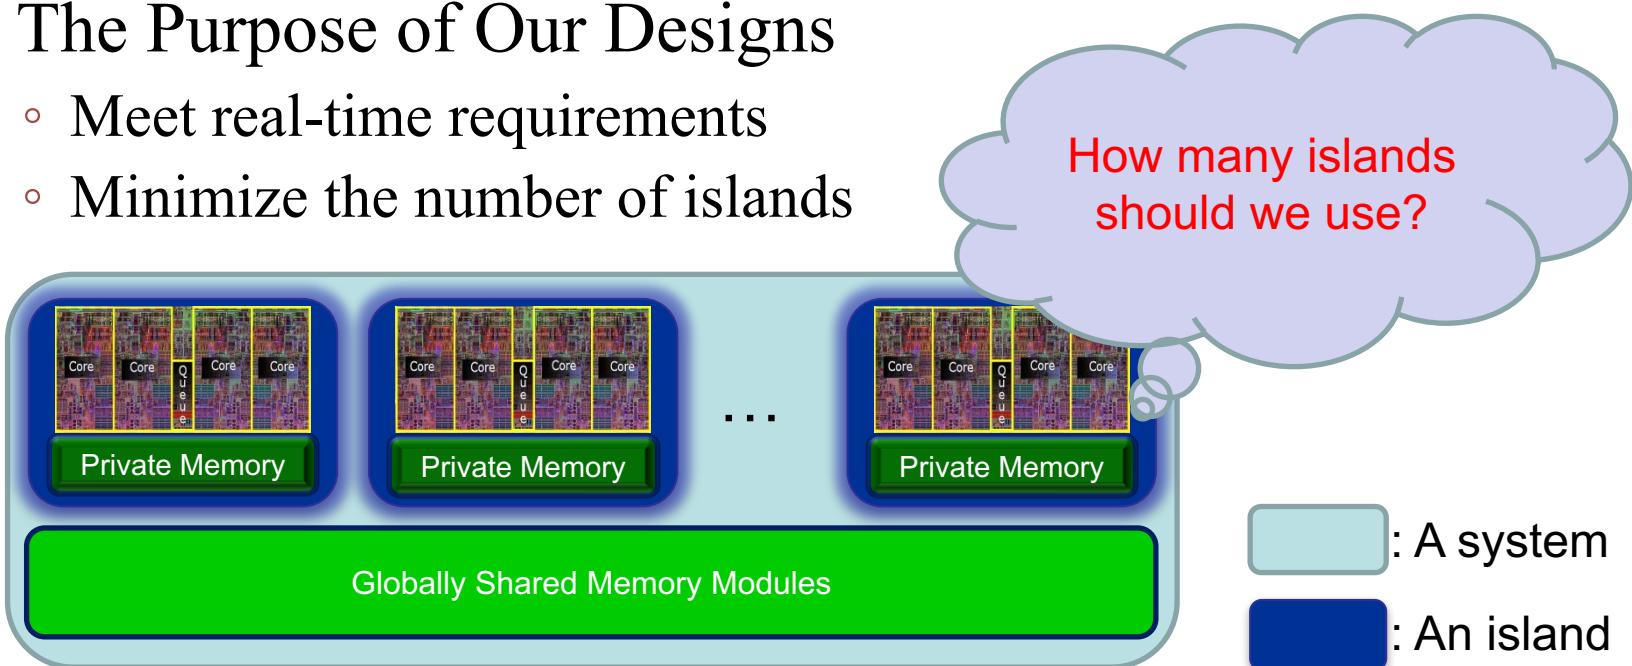
- ▶ The Increasing CPU Performance and Number of Cores
- ▶ The Popularity of Heterogeneous Memory Designs
- ▶ The Conservative Strategies of Traditional Real-Time Scheduling Algorithm Designs



- Source :http://members.tripod.com/pc_tutor1/memory.htm , http://www.infostor.com/index/articles/display/3844378763/articles/infostor/volume-13/Issue_10/Features/Where_do_SSDDs_fit_Servers_or_arrays_.html , <http://www.samsung.com/global/business/semiconductor/product/computing-dram/overview>

System Model

- ▶ An Island-Based Multi-Core Platform
 - A global memory pool and multiple islands
 - A policy to turn on/off islands to manage resources
- ▶ The Purpose of Our Designs
 - Meet real-time requirements
 - Minimize the number of islands

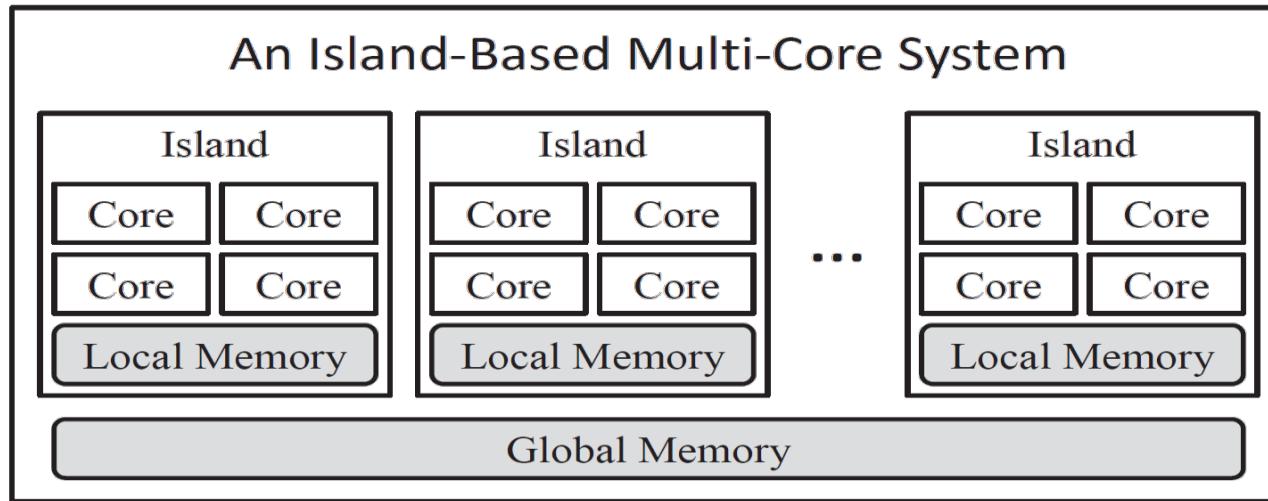


Problem Definition– Real-Time Tasks

- ▶ Goal:
 - Minimize the number of required islands
- ▶ Constraints:
 - All implicit-deadline sporadic tasks meet their deadlines
 - No memory space limitation is violated
- ▶ Task Model— Implicit-Deadline Sporadic Tasks:
 - A relative deadline
 - The minimum inter-arrival time
 - The worst case execution time

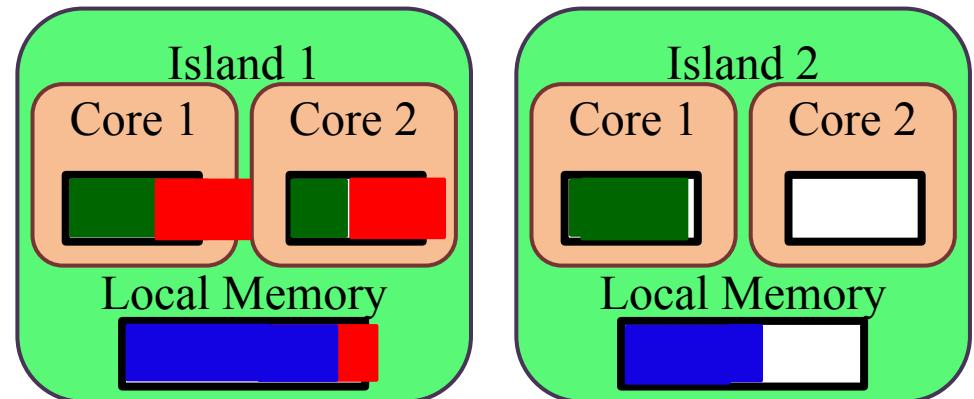
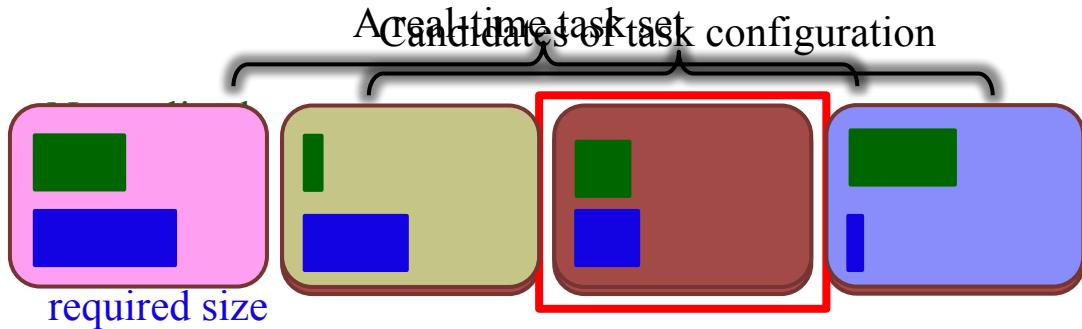
Problem Definition– Hardware Platforms

- ▶ Hardware Model— Island-Based Multi-Core Platforms:
 - Cores in an island share a local memory pool
 - An island consists of multiple homogeneous cores
 - A system consists of multiple islands
 - A system consists of a large global memory pool



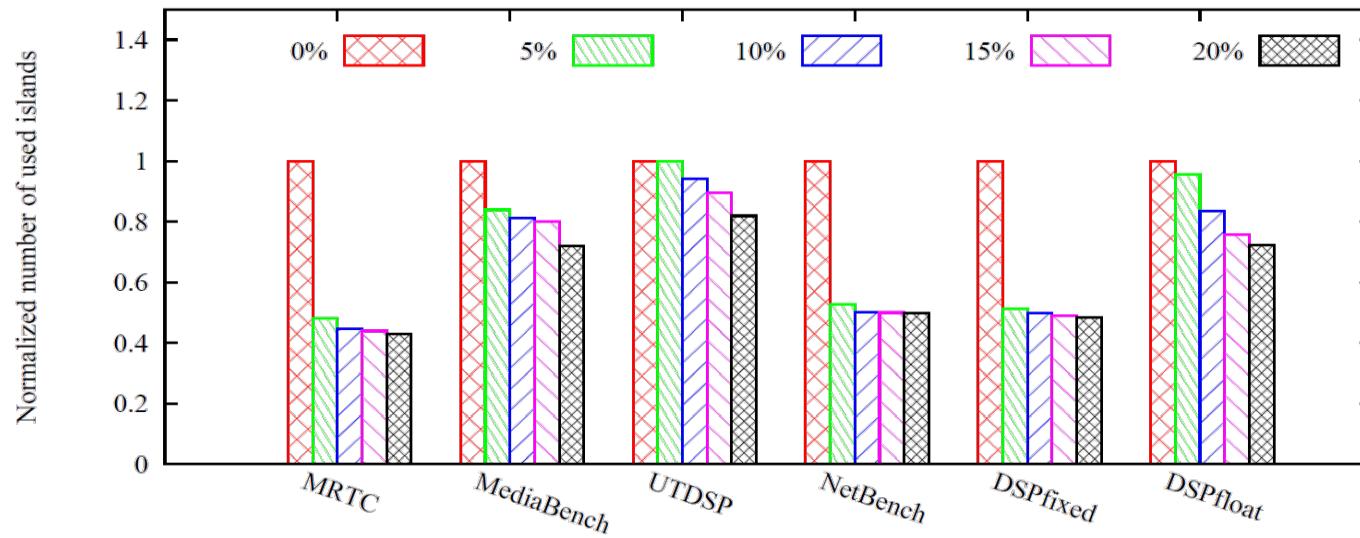
Proposed Algorithms

1. Minimize the value of $\frac{U_i^j}{M} + \frac{j}{B_L}$ for each task
2. Sort all tasks in a non-increasing order by the required number of the local memory blocks
3. Partition all tasks to islands by the first fit strategy
4. Assign tasks onto cores by the first fit strategy
5. Allocate a new island if there is no feasible assignment for a task



Performance Evaluation

- ▶ 82 real-life benchmarks from MRTC, MediaBench, UPDSP, NetBench, and DSPstone
- ▶ The worst-case execution time of each task is generated by aiT which is based on Infineon TriCore TC1797
- ▶ The number of cores in an island is 4



Conclusion

▶ Summary

- Minimize the needed number of islands
- Meet the real-time required
- Provide a polynomial-time asymptotic 4-approximation algorithm

▶ Insights

- The design of the memory architecture of multi-core systems is with increasing importance
- CPU is not the only resource which should be considered in multi-core real-time scheduling



Summarizing this Semester

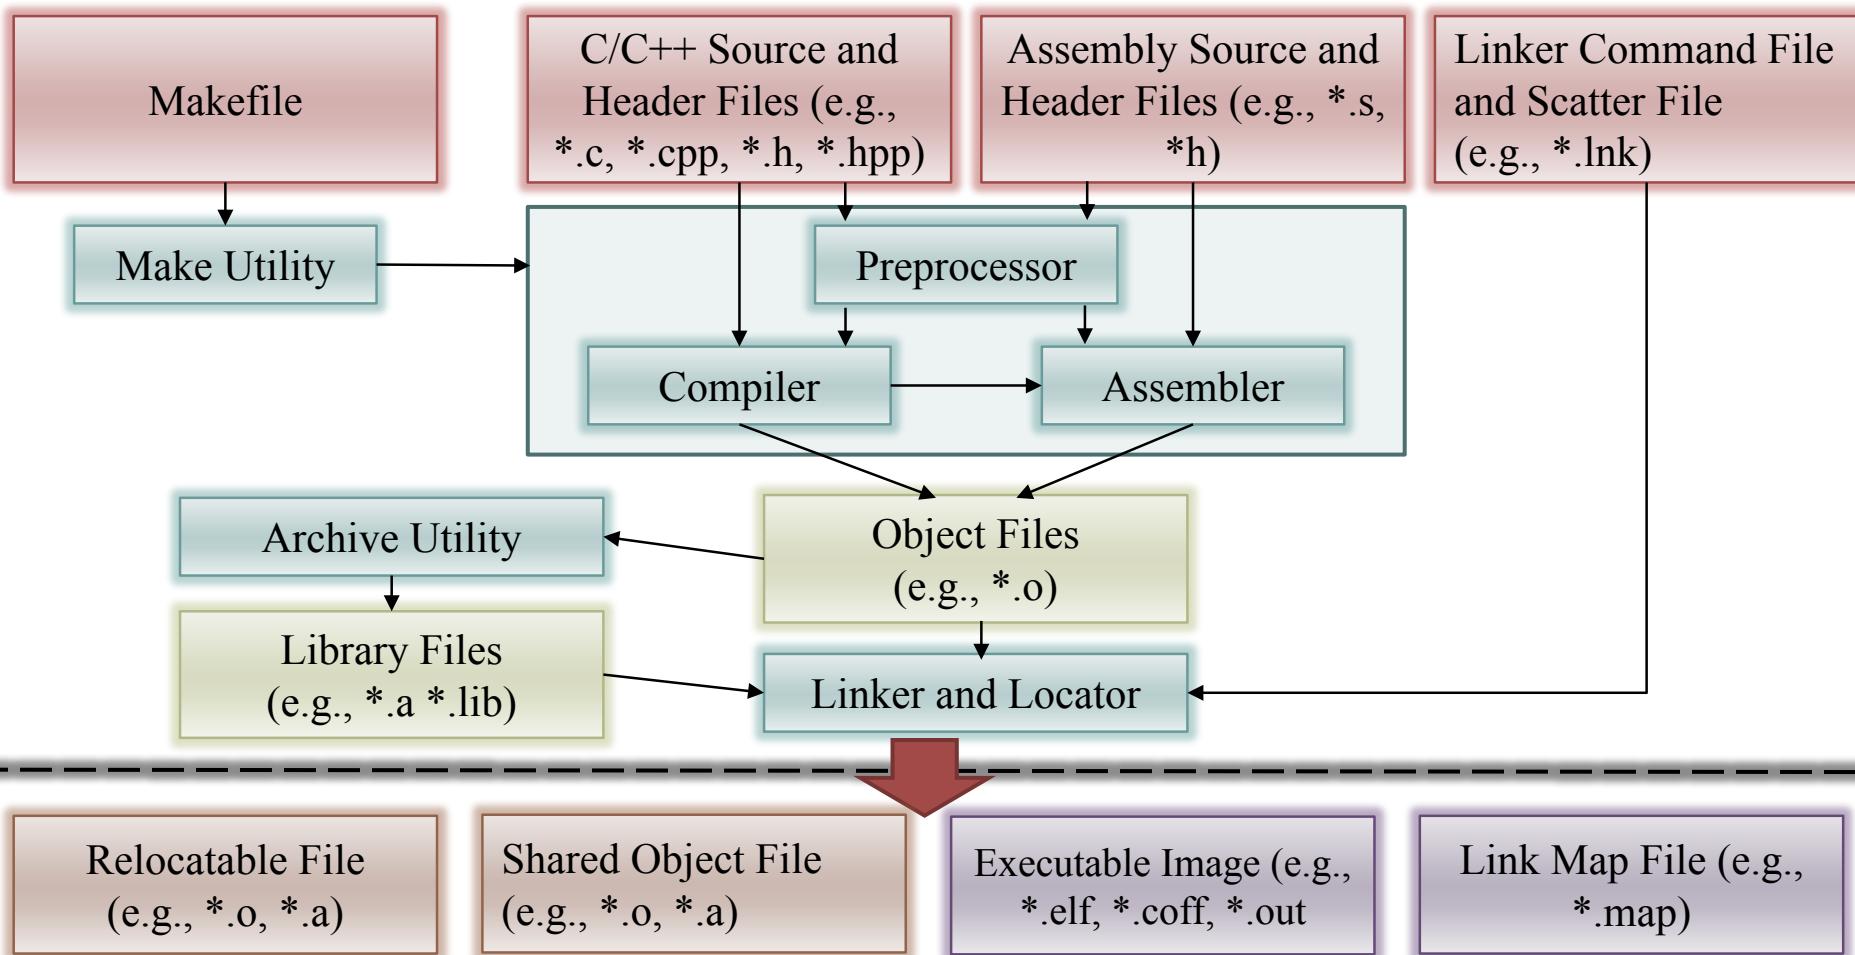
Common Performance Metrics

- ▶ **Unit Cost**: the monetary cost of manufacturing each copy of the system
- ▶ **NRE Cost** (Non-Recurring Engineering cost): the one-time monetary cost of designing the system
- ▶ **Size**: the physical space required by the system
- ▶ **Performance**: the execution time or throughput of the system
- ▶ **Power**: the amount of power consumed by the system
- ▶ **Flexibility**: the ability to change the functionality of the system without incurring heavy NRE cost
- ▶ **Time-to-Market**: the time required to develop a system to the point that it can be released and sold to customers
- ▶ **Maintainability**: the ability to modify the system after its initial release

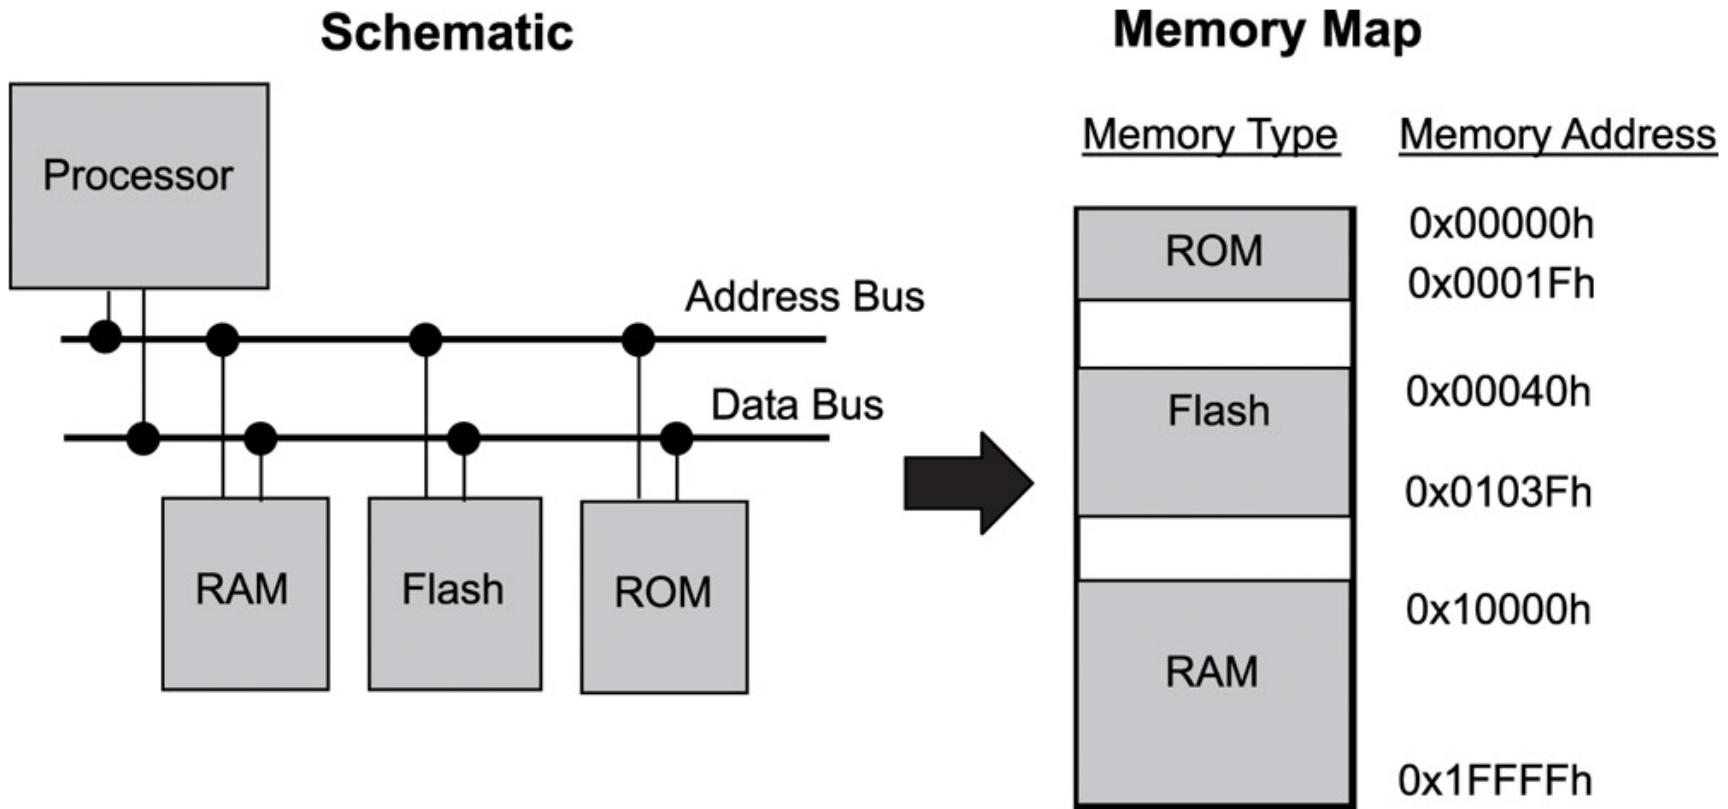
Definition of Embedded Systems

- ▶ It is difficult to define embedded systems
 - An embedded system is a digital system
 - An embedded system has computing processors
 - An embedded system runs dedicated functions
 - An embedded system is frequently used as a controller
- ▶ An embedded system is a **computer system** with some **dedicated functions** within a larger **mechanical or electrical system**, often with **real-time** computing constraints. - From Wikipedia

Creating an Image for a Target Embedded System



Simplified Schematic and Memory Map for a Target System



Mapping an Executable Image into the Target System

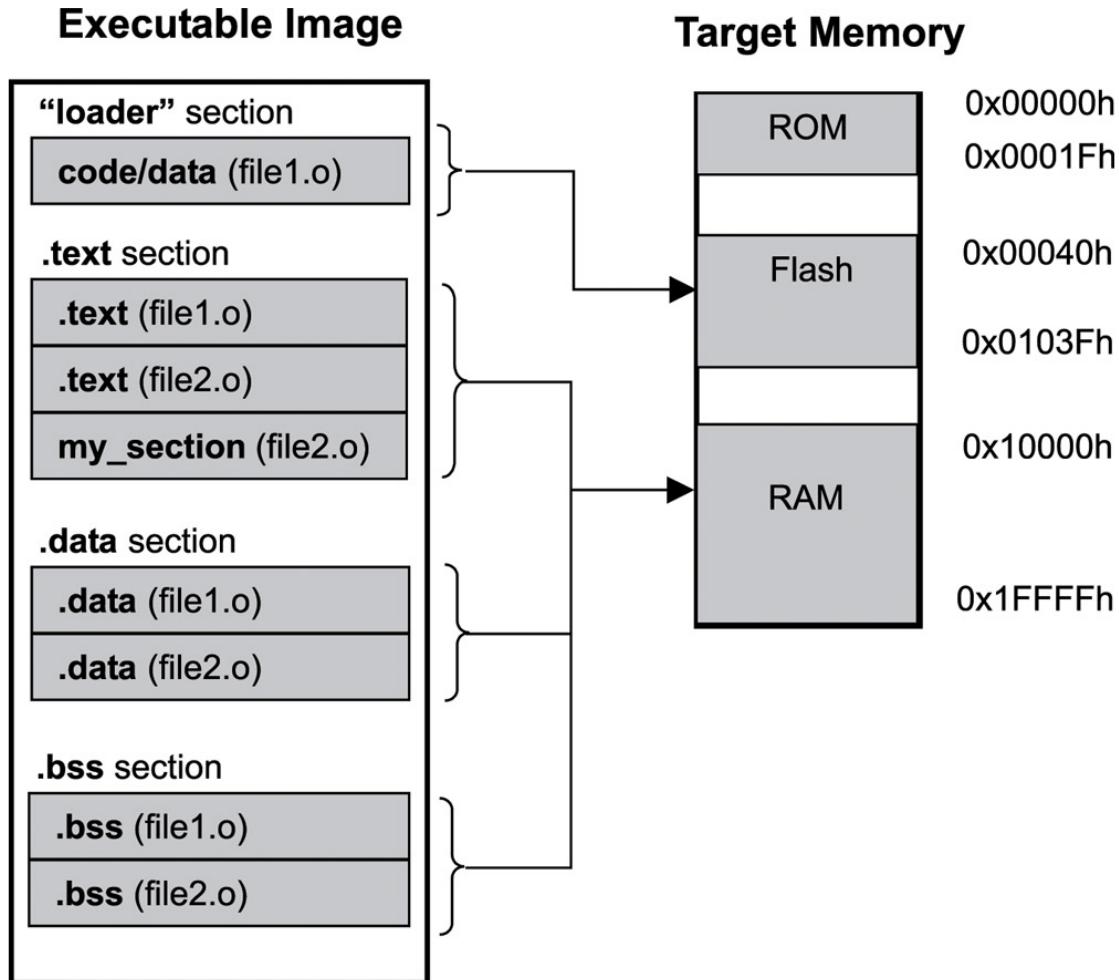


Image Transferred from ROM Running on RAM (1/2)

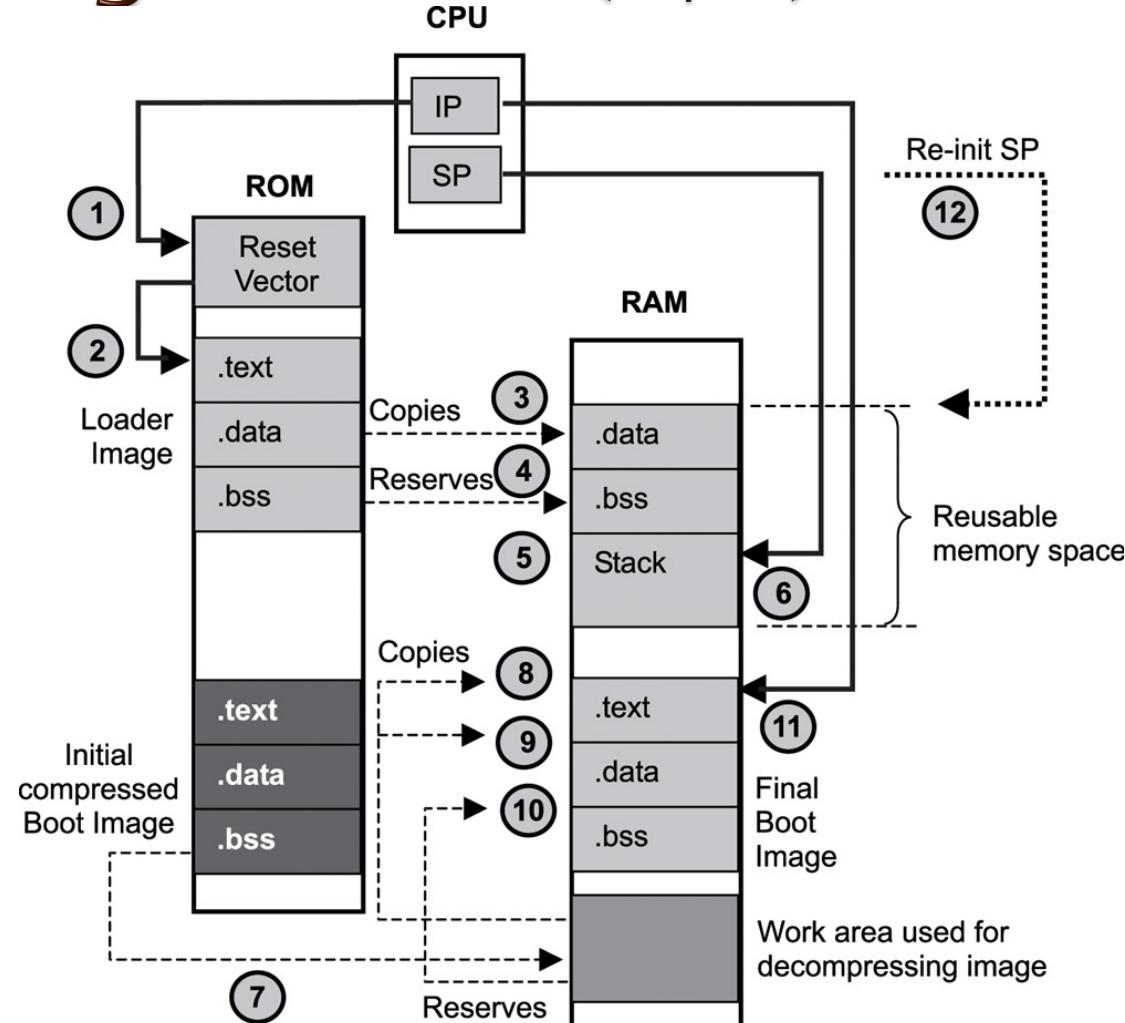


Image Transferred from ROM Running on RAM (2/2)

- ▶ The boot loader transfers an application image from ROM to RAM for execution
 - The application image is usually compressed in ROM to reduce the storage space required
- ▶ Boot sequence
 7. Copy the Compressed application image from ROM to RAM in a work area
 8. Decompress and initialize the application image(1)
 9. Decompress and initialize the application image(2)
 10. Decompress and initialize the application image(3)
 11. The loader transfers control to the image using a processor-specific jump instruction
 12. Recycle the memory area occupied by the loader and the work area
 - May also reinitialize the SP to point to the memory area occupied by the loader to use it as the stack space

An Example of Real-Time Designs

- ▶ A camera periodically takes a photo
- ▶ The image recognition result will be produced before the next period
- ▶ If there is an obstacle, the train automatically brakes

$$\text{Time of a Period} = 150/50 = 3\text{s}$$

$$\text{Distance of a Period} = (400 - 100)/2 = 150\text{m}$$

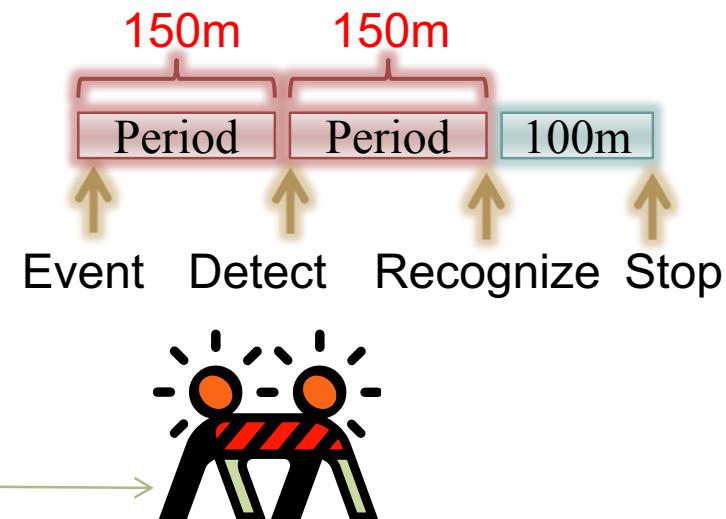
Braking: -12.5m/s^2

Max Speed: 50m/s



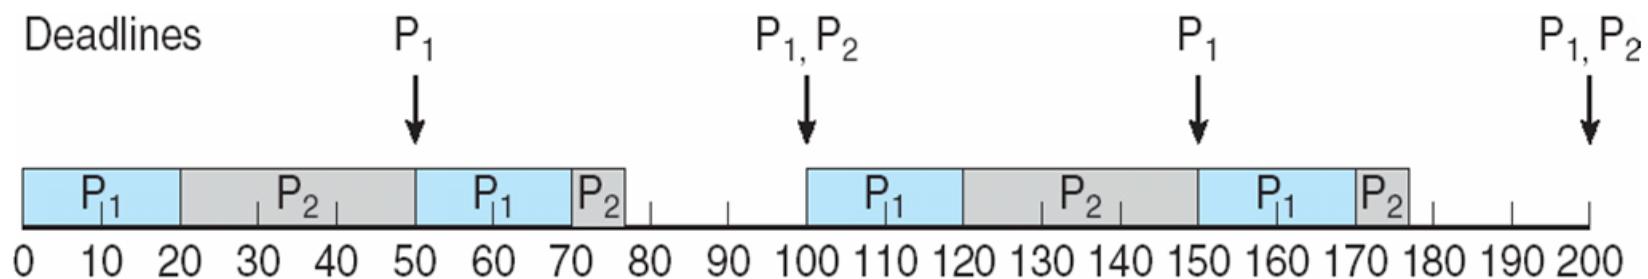
$$\text{Distance to Stop} \\ 25 \times (50/12.5) = 100\text{m}$$

Camera Range: 400m



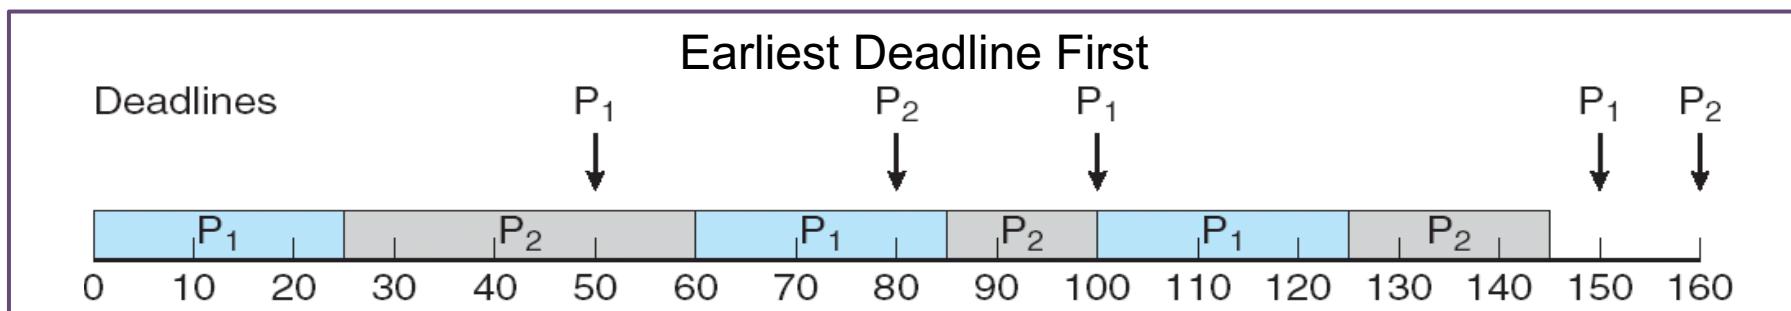
A Static Scheduling Algorithm— Rate Monotonic Scheduling

- ▶ A static priority is assigned to each task based on the inverse of its period
 - A task with shorter period → higher priority
 - A task with longer period → lower priority
 - For example:
 - P_1 has its **period 50** and execution time 20
 - P_2 has its **period 100** and execution time 37
 - P_1 is assigned a higher priority than P_2



A Dynamic Scheduling Algorithm—Earliest Deadline First Scheduling

- ▶ Dynamic priorities are assigned according to deadlines
 - The earlier the deadline, the higher the priority
 - The later the deadline, the lower the priority
 - For example:
 - P_1 has its period 50 and execution time 25
 - P_2 has its period 80 and execution time 35



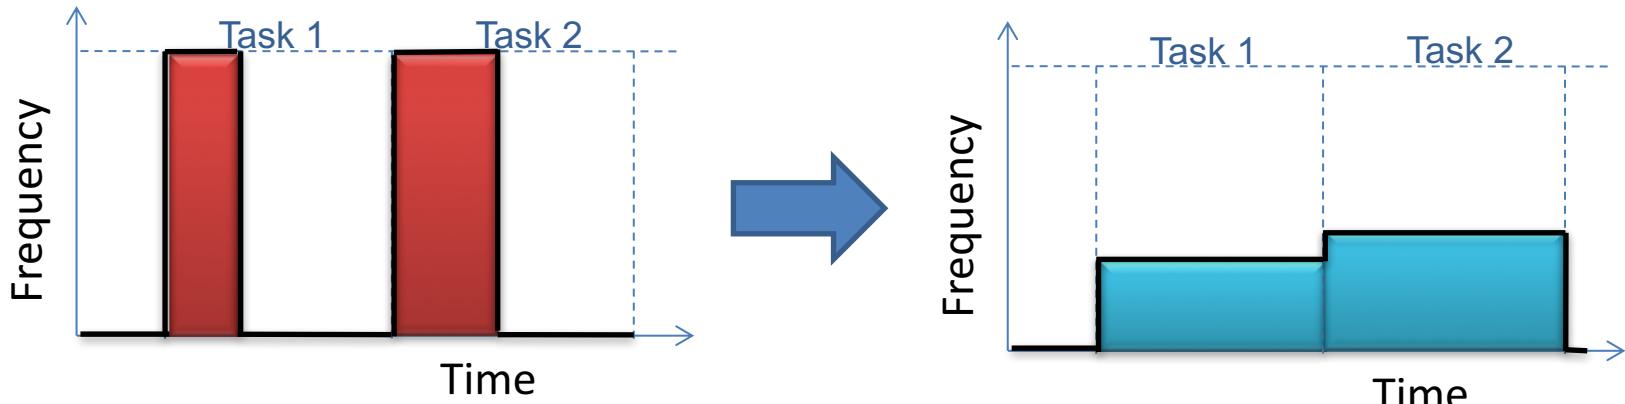
Two Approaches to Reduce the Power Consumption of Devices

- ▶ Dynamic Voltage and Frequency Scaling (DVFS)
 - Scale down the voltage and/or frequency to reduce the processor power consumption
 - An example: reducing 17% of the energy consumption for H.264 decoding over TI DaVince DM6446
- ▶ Dynamic Power Management (DPM)
 - Change to an energy-efficient state to reduce the power consumption of peripheral devices
 - An example: reducing 15% of the energy consumption for the browser over Android Galaxy Tab

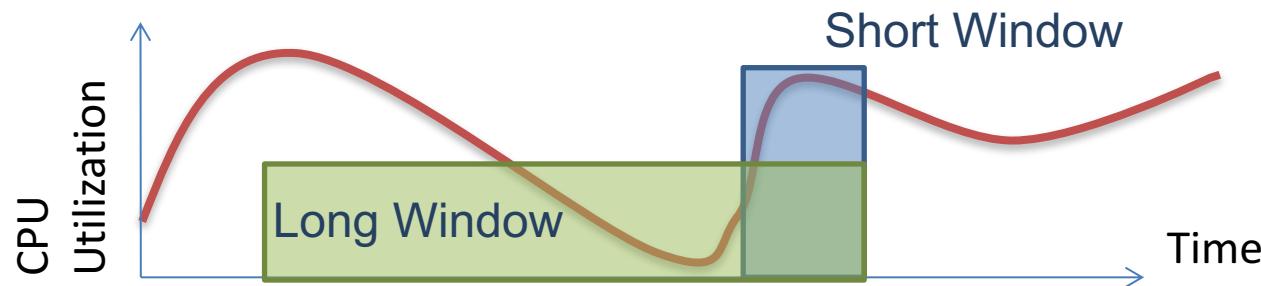


Dynamic Frequency and Voltage Scaling Designs

- The Observation: It is energy-efficient to scale down the processor frequency dynamically to fit current workloads



- The Approach: It keeps a window or buffer to estimate the workload



Introduction of μC/OS-II (1 / 2)

- ▶ The name is from micro-controller operating system, version 2
- ▶ μC/OS-II is certified in an avionics product by FAA in July 2000 and is also used in the Mars Curiosity Rover
- ▶ It is a very small real-time kernel
 - Memory footprint is about 20KB for a fully functional kernel
 - Source code is about 5,500 lines, mostly in ANSI C
 - Its source is open but not free for commercial usages
- ▶ Preemptible priority-driven real-time scheduling
 - 64 priority levels (max 64 tasks)
 - 8 reserved for μC/OS-II
 - Each task is an infinite loop



Introduction of μC/OS-II (2/2)

- ▶ Deterministic execution times for most μC/OS-II functions and services
- ▶ Nested interrupts could go up to 256 levels
- ▶ Supports of various 8-bit to 64-bit platforms: x86, 68x, MIPS, 8051, etc.
- ▶ Easy for development: Borland C++ compiler and DOS (optional)
- ▶ However, uC/OS-II still lacks of the following features:
 - Resource synchronization protocol
 - Soft-real-time support

The μC/OS-II File Structure

Application Code (Your Code!)

Processor Independent Implementations

- Scheduling policy
- Event flags
- Semaphores
- Mailboxes
- Event queues
- Task management
- Time management
- Memory management

Application Specific Configurations

- OS_CFG.H
- Max # of tasks
- Max Queue length
- ...

uC/OS-II Port for Processor Specific Codes

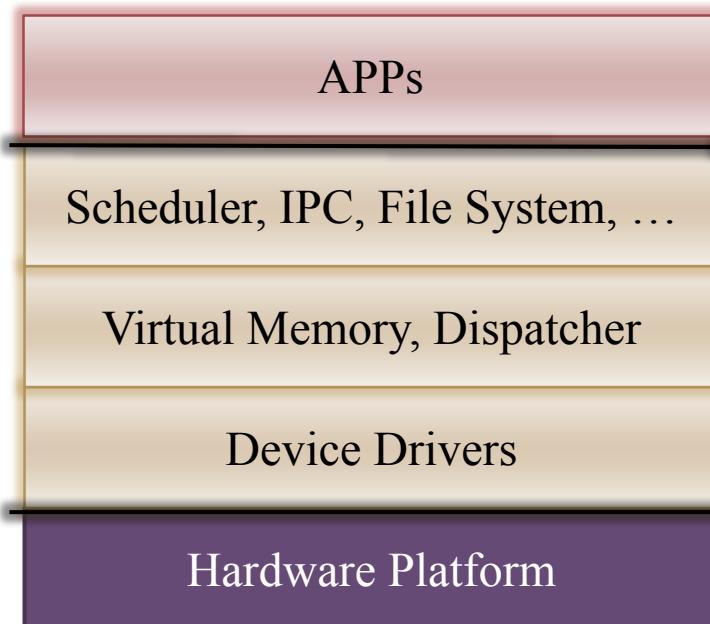
Software
Hardware

CPU

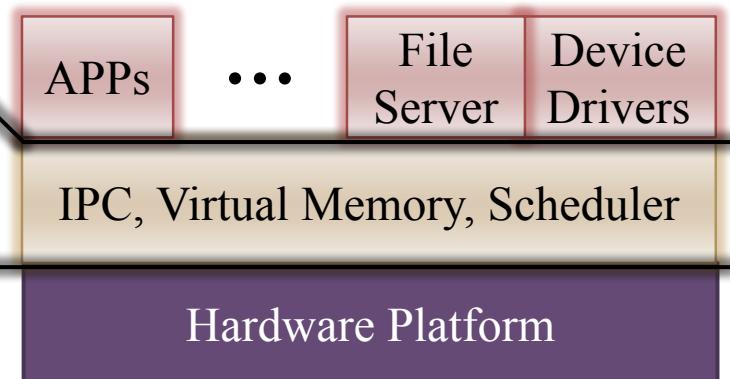
Timer

Monolithic Kernel and Microkernel

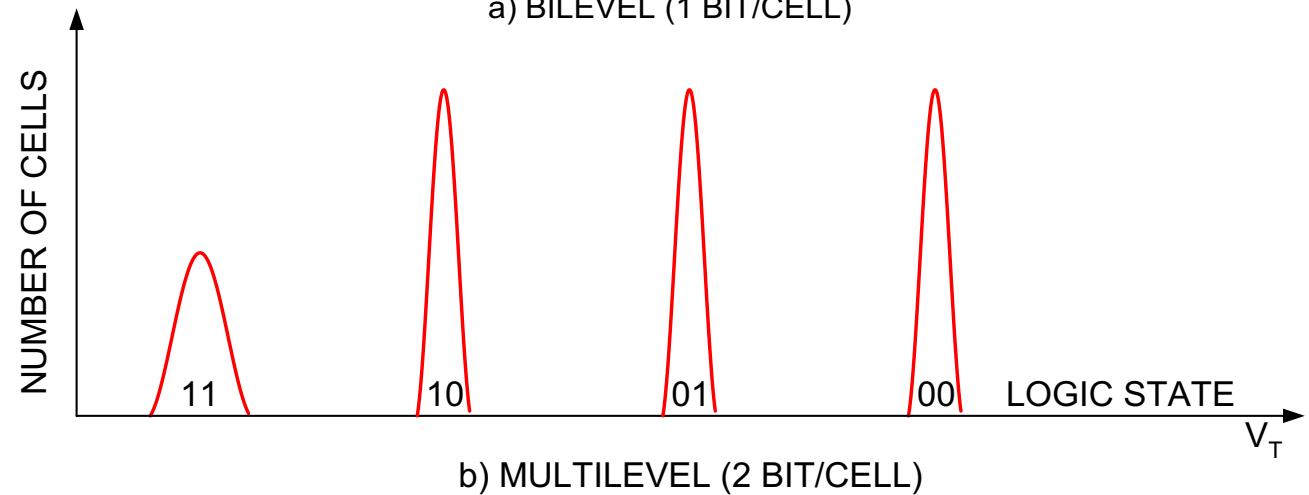
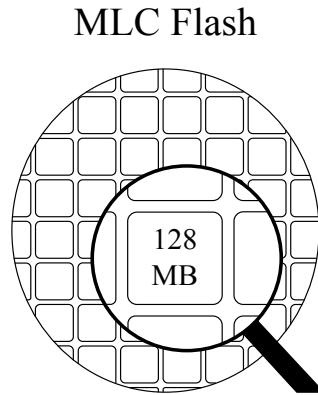
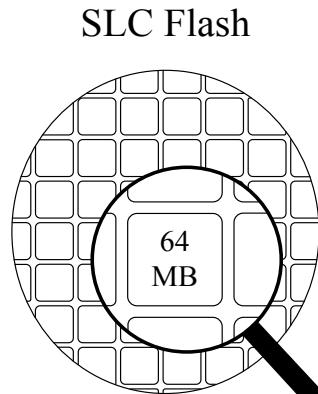
Monolithic Kernel



Microkernel



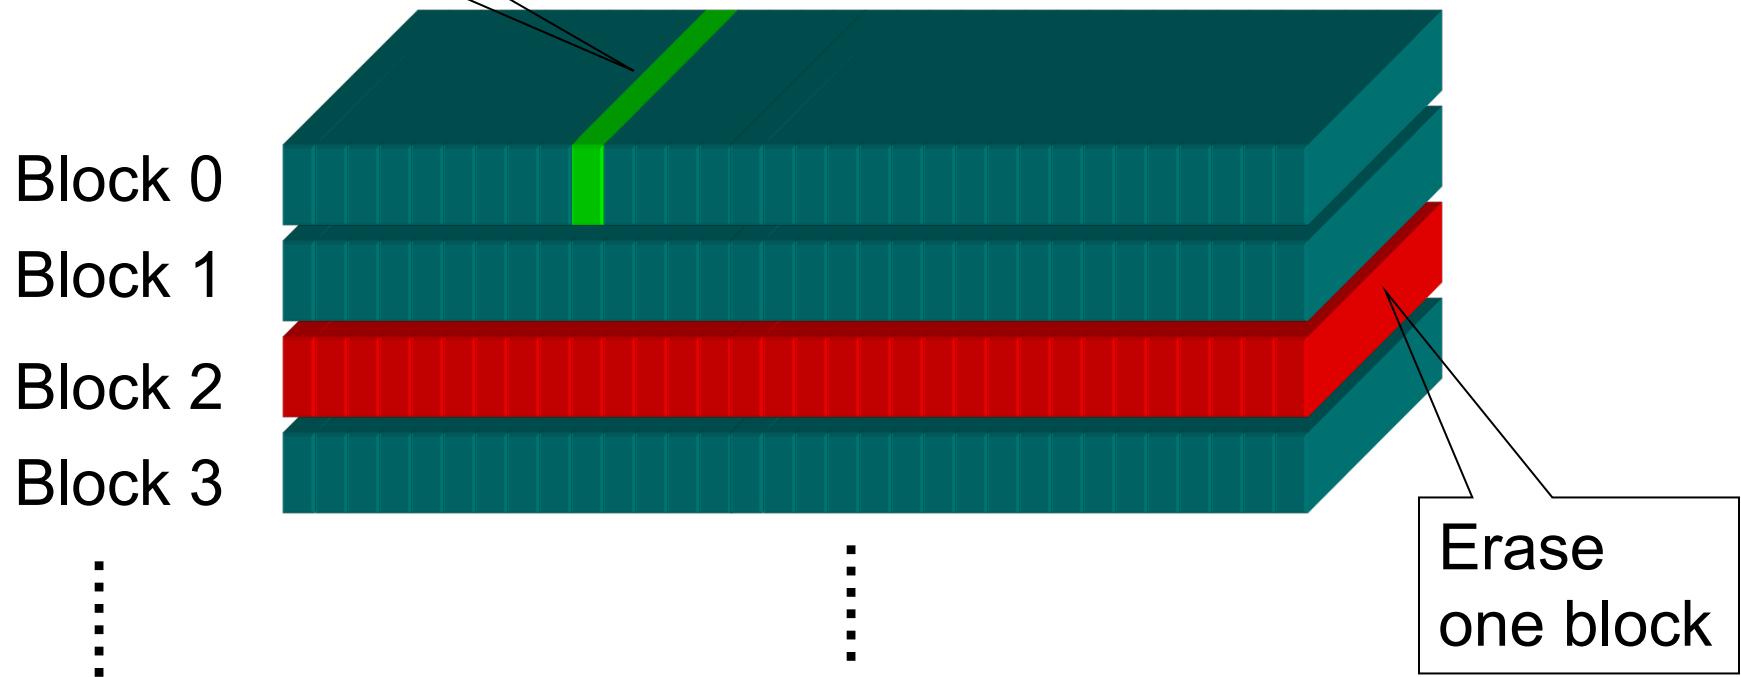
Single-Level Cell (SLC) vs Multi-Level Cell (MLC) Flash



Page Write and Block Erase

Write one page

1 Page = 512B—4KB
1 Block = 32 pages



Lab 1: Android Tools

- ▶ Goal: run android applications on Qualcomm evaluation boards
- ▶ Install JDK
- ▶ Install Android Studio
- ▶ Create a project and use API level 19 (for kitkat)
- ▶ Run the program on an Android emulator
- ▶ Export the apk file and upload it to somewhere
- ▶ Download it to the evaluation board and run it

Lab 2: Embedded System Tools and Multi-thread Programs

- ▶ Build the Cross Development Toolchain
- ▶ Build the Linux Kernel
- ▶ Setup a TFTP Server
- ▶ Setup NFS Server
- ▶ Setup the Target Board
- ▶ Download the Linux Kernel
- ▶ Write Multi-thread programs on TI OMAP Evaluation Board
- ▶ Use Mutex to protect your program



Thank You!