



Operating System Concepts

Che-Wei Chang

chewei@mail.cgu.edu.tw

Department of Computer Science and Information
Engineering, Chang Gung University

Contents



1. Introduction
2. System Structures
3. Process Concept
4. Multithreaded Programming
5. Process Scheduling
6. Synchronization
7. Deadlocks
8. Memory-Management Strategies
9. Virtual-Memory Management
10. File System
11. Implementing File Systems
12. Secondary-Storage Systems

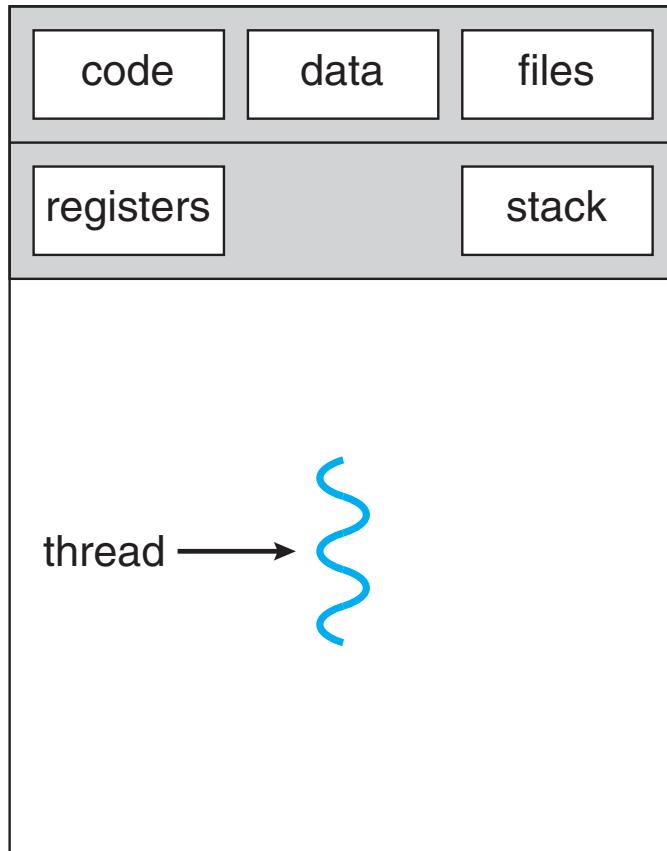


Chapter 4. Multithreaded Programming

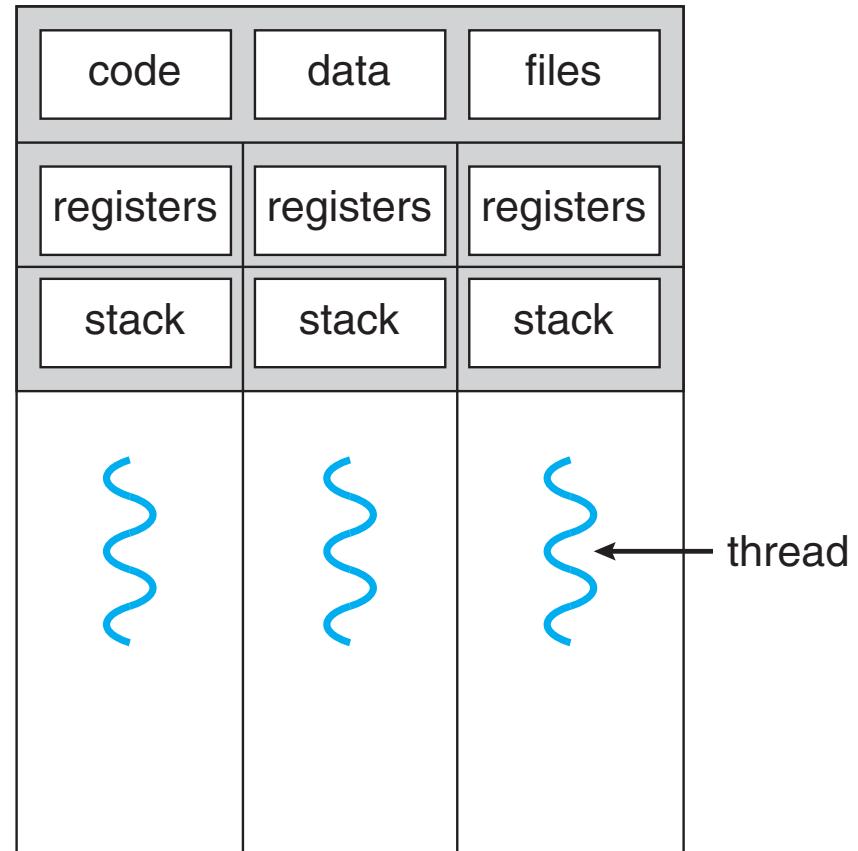
Objectives

- ▶ To introduce the notion of a thread
- ▶ To discuss the APIs for the Pthreads, Windows, and Java thread libraries
- ▶ To explore several strategies that provide implicit threading
- ▶ To examine issues related to multithreaded programming

Single and Multithreaded Processes

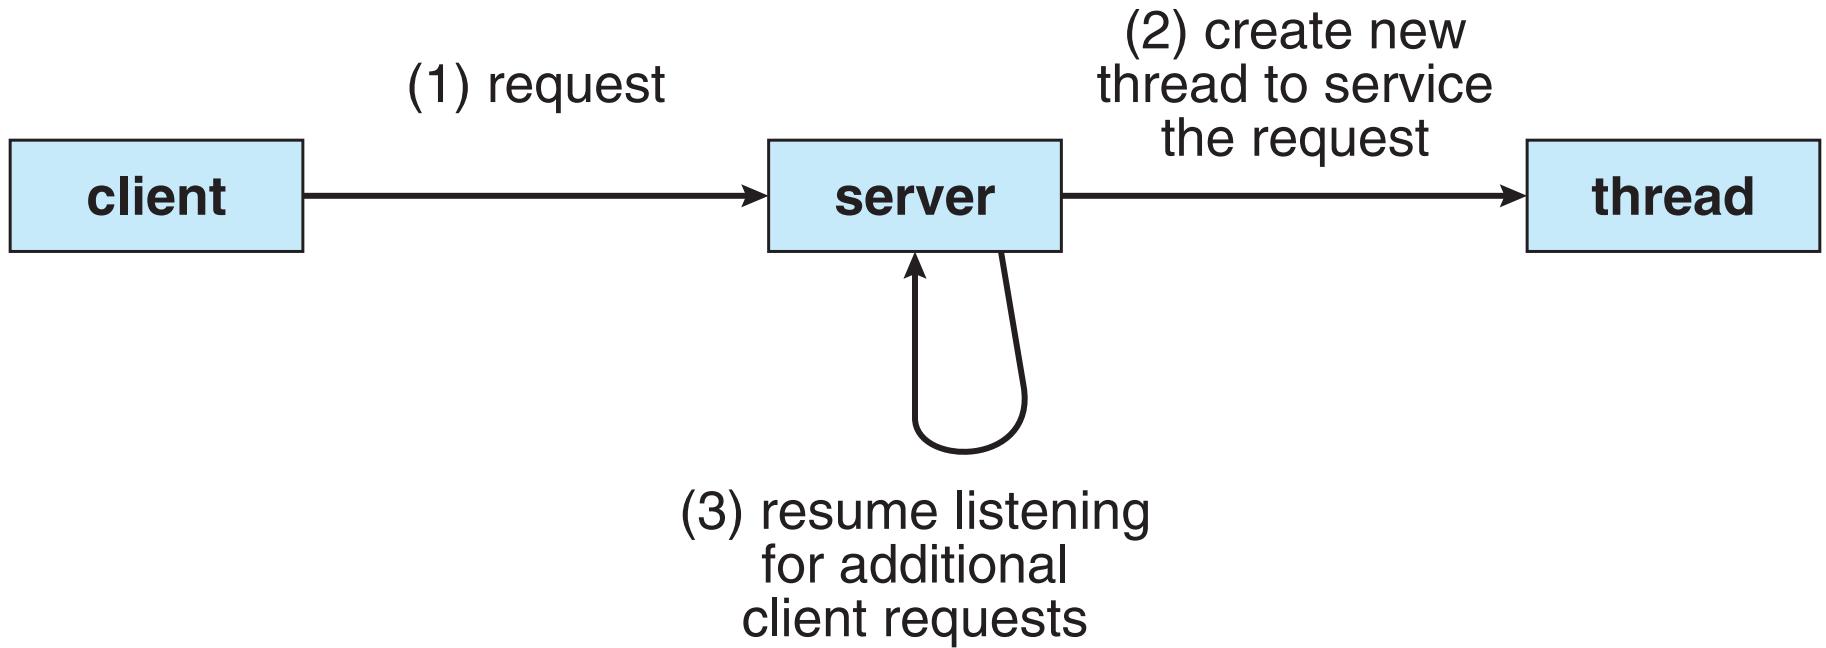


single-threaded process



multithreaded process

Multithreaded Server Architecture



Motivation

- ▶ Most modern applications are multithreaded
- ▶ Multiple tasks with the application can be implemented by separate threads
 - Update display
 - Fetch data
 - Spell checking
- ▶ Process creation is heavy-weight while thread creation is light-weight
- ▶ Kernels are generally multithreaded

Benefits

- ▶ Responsiveness
 - It allows a program to continue running even if part of it is blocked or is performing a lengthy operation
- ▶ Resource Sharing
 - Threads share resources of process, easier than shared memory or message passing
- ▶ Economy
 - Thread creation is cheaper than process creation
 - Thread switching overhead is lower than context switching
- ▶ Scalability
 - Threads can efficiently use multiprocessor architectures

Multicore Programming

- ▶ Motivation: the popularity of multiple computing cores per system
 - Multithreaded Programming
- ▶ Challenges in Programming
 - Dividing Activities
 - Load Balancing
 - Data Splitting
 - Data Dependency
 - Testing and Debugging

User Threads and Kernel Threads

▶ User threads

- Management done by user-level threads library
- Three primary thread libraries:
 - POSIX Pthreads
 - Win32 threads
 - Java threads

▶ Kernel threads

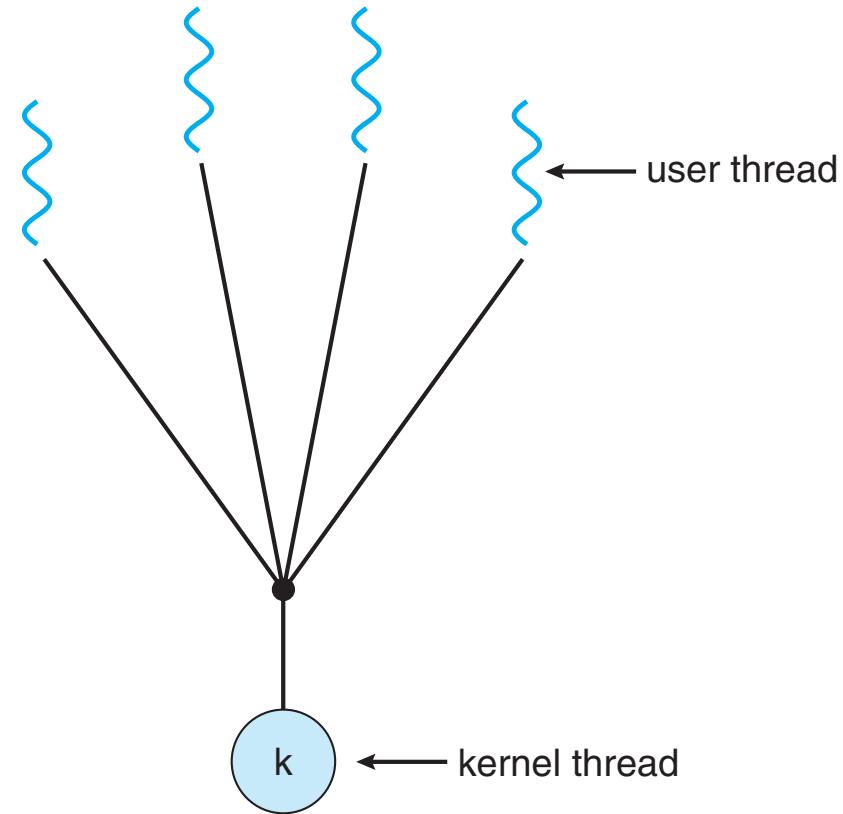
- Supported by the Kernel
- Examples – virtually all general purpose operating systems, including:
 - Windows, Solaris, Linux, Tru64 UNIX, Mac OS X

Multithreading Models

- ▶ Relationship between user threads and kernel threads
 - Many-to-One
 - One-to-One
 - Many-to-Many

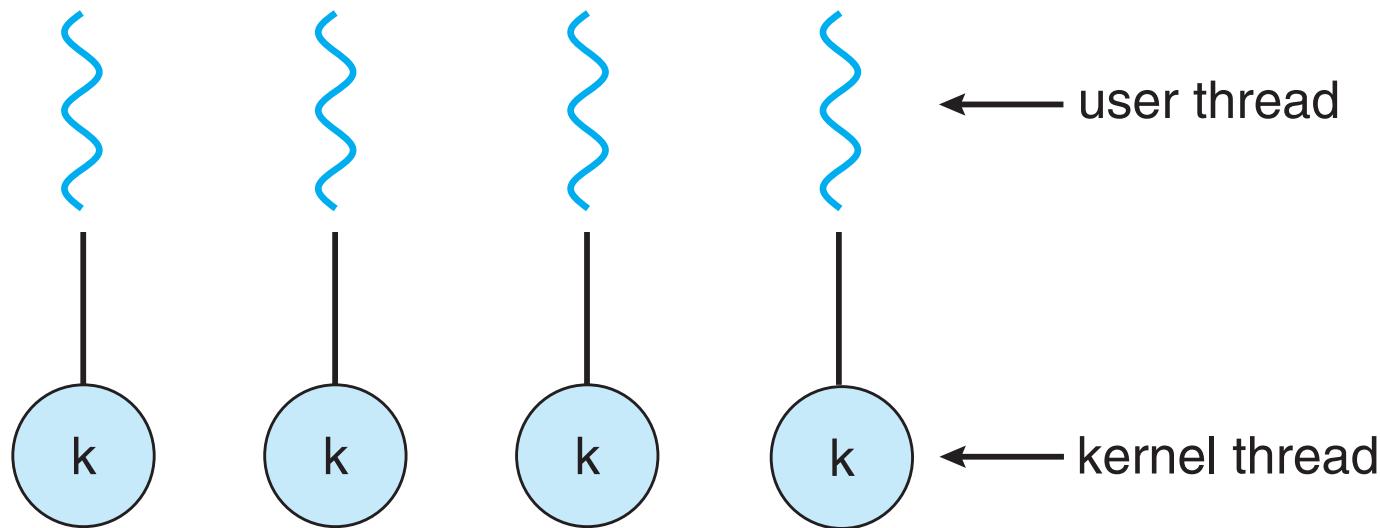
Many-to-One Model

- ▶ Many user threads to one kernel thread
- ▶ Advantage:
 - Efficiency
- ▶ Disadvantage:
 - One blocking system call blocks all
 - No parallelism for multiple processors
- ▶ Example:
 - Solaris Green Threads
 - GNU Portable Threads



One-to-One Model

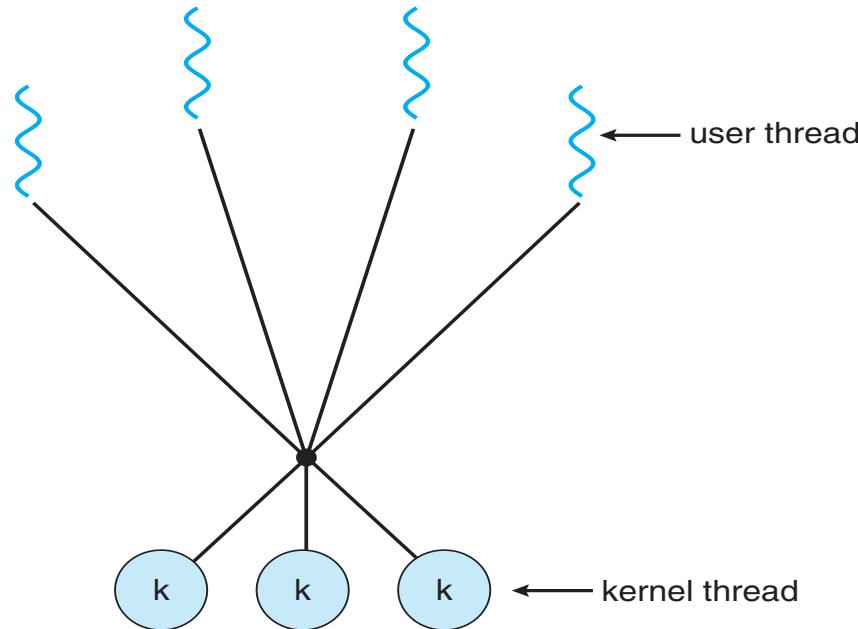
- ▶ One user-level thread to one kernel thread
- ▶ Advantage: One system call blocks one thread
- ▶ Disadvantage: Overheads in creating a kernel thread
- ▶ Example: Windows NT/2000/XP, Linux, Solaris 9



Many-to-Many Model

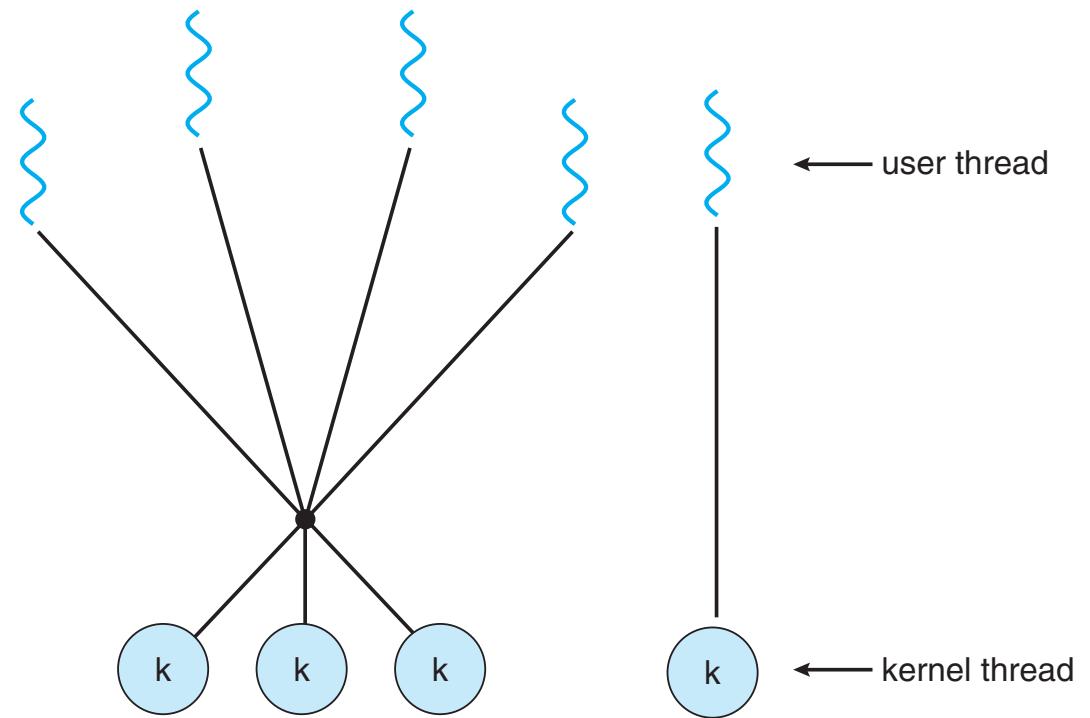
▶ Many-to-Many Model

- Many user-level threads to many kernel threads
- Advantage: A combination of parallelism and efficiency
- Example: Solaris prior to version 9



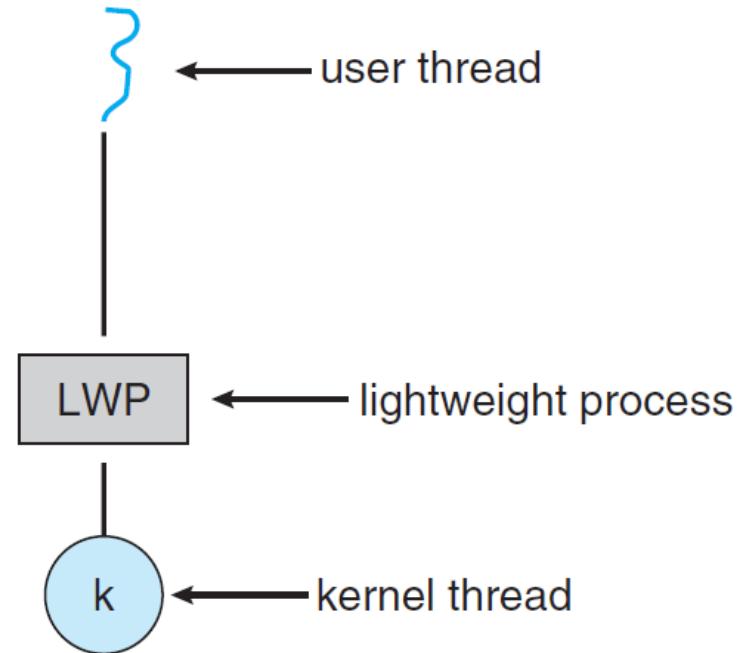
Two-Level Model

- ▶ Similar to the many-to-many model, except that it allows a user thread to be bound to kernel thread
- ▶ Examples
 - IRIX
 - HP-UX
 - Tru64 UNIX
 - Solaris 8 and earlier



Scheduler Activations

- ▶ Definition: A scheme for the communication between the user-thread library and the kernel
 - The kernel provides a set of virtual processors, i.e., light weight processes (LWP)
 - User threads on a LWP are blocked if any of the user threads is blocked!



Thread Libraries

- ▶ The goal thread libraries is to provide an API for creating and managing threads
- ▶ Two Approaches
 - User Thread Library
 - Kernel-Level Thread Library
- ▶ Well-Known Examples
 - POSIX Pthread – User or Kernel Level
 - Win32 thread – Kernel Level
 - Java thread – Level Depending on the Thread Library on the Host System

A Pthread Example (1/3)

- ▶ The specification of the example program
 - Read an input integer N
 - Create a thread to calculate the summation from 1 to N
 - Wait for the completion of the thread
 - Print the result from the thread
- ▶ Now, let's use the Pthread library to implement the program

A Pthread Example (2/3)

```
#include <pthread.h>
#include <stdio.h>

int sum; /* this data is shared by the thread(s) */
void *runner(void *param); /* threads call this function */

int main(int argc, char *argv[])
{
    pthread_t tid; /* the thread identifier */
    pthread_attr_t attr; /* set of thread attributes */

    if (argc != 2) {
        fprintf(stderr,"usage: a.out <integer value>\n");
        return -1;
    }
    if (atoi(argv[1]) < 0) {
        fprintf(stderr,"%d must be >= 0\n",atoi(argv[1]));
        return -1;
    }
}
```

Include the header file of pthread

Declare the function to be executed by the thread

Create the data-structure to be used by the thread

A Pthread Example (3/3)

```
/* get the default attributes */
pthread_attr_init(&attr);
/* create the thread */
pthread_create(&tid,&attr,runner,argv[1]);
/* wait for the thread to exit */
pthread_join(tid,NULL);
printf("sum = %d\n",sum);
}

/* The thread will begin control in this function */
void *runner(void *param)
{
    int i, upper = atoi(param);
    sum = 0;

    for (i = 1; i <= upper; i++)
        sum += i;

    pthread_exit(0);
}
```

Initialize the data-structure to be used by the thread

Create the thread

Wait for the completion of the thread

Define the function to be executed by the thread

Compiling POSIX-Thread Programs

Compiler / Platform	Compiler Command	Description
INTEL Linux	icc -pthread	C
PGI Linux	icpc -pthread	C++
GNU Linux, Blue Gene	pgcc -lpthread pgCC -lpthread	C
GNU Linux, Blue Gene	gcc -lpthread g++ -lpthread	C C++
IBM Blue Gene	bgxlC_r / bgcc_r bgxlC_r, bgxlC++_r	GNU C++ C (ANSI / non-ANSI) C++

Implicit Threading

- ▶ Implicit threading is growing in popularity as numbers of threads increase
- ▶ Program correctness is more difficult with explicit threads
- ▶ Creation and management of threads done by compilers and run-time libraries rather than programmers
- ▶ Examples
 - OpenMP on Linux, Windows and Mac OS X
 - Grand Central Dispatch on Mac OS X

Threading Issues

- ▶ Semantics of **fork()** and **exec()** system calls
- ▶ Signal handling
 - Synchronous and asynchronous
- ▶ Thread cancellation of target thread
 - Asynchronous or deferred
- ▶ Thread-local storage
- ▶ Scheduler activations

Fork and Exec System Calls

- ▶ When a process consists of multiple threads, does **fork()** duplicate only the calling thread or all threads?
 - Some UNIX systems have two versions of fork()
- ▶ **exec()** usually works as normal—replaces the running process including all threads

Signal Handling

- ▶ Two Types of Signals
 - Synchronous signal— should be delivered to the same process that performed the operation causing the signal
 - e.g., illegal memory access or division by zero
 - Asynchronous signal— can happen at any time point
 - e.g., ^C or timer expiration
- ▶ Delivery of a Signal
 - To the thread to which the signal applies
 - e.g., division-by-zero
 - To every thread in the process
 - e.g., ^C
 - To certain threads in the process
 - Assign a specific thread to receive all signals for the process

Thread Cancellation

- ▶ A cancellation signal is sent to the target thread
- ▶ Two scenarios for the cancellation:
 - Asynchronous cancellation
 - Immediate cancel the thread
 - Deferred cancellation
 - Wait until some special point of the thread, e.g., cancellation points in Pthread
- ▶ Difficulty
 - Resources have been allocated to a cancelled thread
 - A thread is cancelled while it is updating data

Thread-Local Storage

- ▶ Thread-local storage (TLS) allows each thread to have its own copy of data
- ▶ Different from local variables
 - Local variables visible only during single function invocation
 - TLS visible across function invocations
- ▶ Similar to **static** data
 - TLS is unique to each thread

Windows Threads

- ▶ Windows implements the Windows API— primary API for Win 98, Win NT, Win 2000, Win XP, Win 7, Win 8, and Win 10
- ▶ It implements the one-to-one mapping
- ▶ Each thread contains
 - A thread id
 - Register set representing state of processor
 - Separate user and kernel stacks for when thread runs in user mode or kernel mode
 - Private data storage area used by run-time libraries and dynamic link libraries (DLLs)
- ▶ The register set, stacks, and private storage area are known as the **context** of a thread

Linux Threads

- ▶ The concepts of threads was introduced in version 2.2
- ▶ In Linux
 - Processes and threads are called tasks
 - Any task has a PID (process identifier)
 - If two tasks do not share any data-structure, they are two processes
 - If two tasks share some data-structure, they just like two threads in the same process
 - `fork()` is used to create a new process
 - `clone()` is used to create a new thread
 - Flag setting in `clone()` invocation: `CLONE_FS`, `CLONE_VM`, `CLONE_SIGHAND`, `CLONE_FILES`