



Operating System Concepts

Che-Wei Chang

chewei@mail.cgu.edu.tw

Department of Computer Science and Information
Engineering, Chang Gung University

Contents

1. Introduction
2. System Structures
3. Process Concept
4. Multithreaded Programming
5. Process Scheduling
6. Synchronization
7. Deadlocks
8. Memory-Management Strategies
9. Virtual-Memory Management
10. File System
11. Implementing File Systems
12. Secondary-Storage Systems



Chapter 3. Process Concept

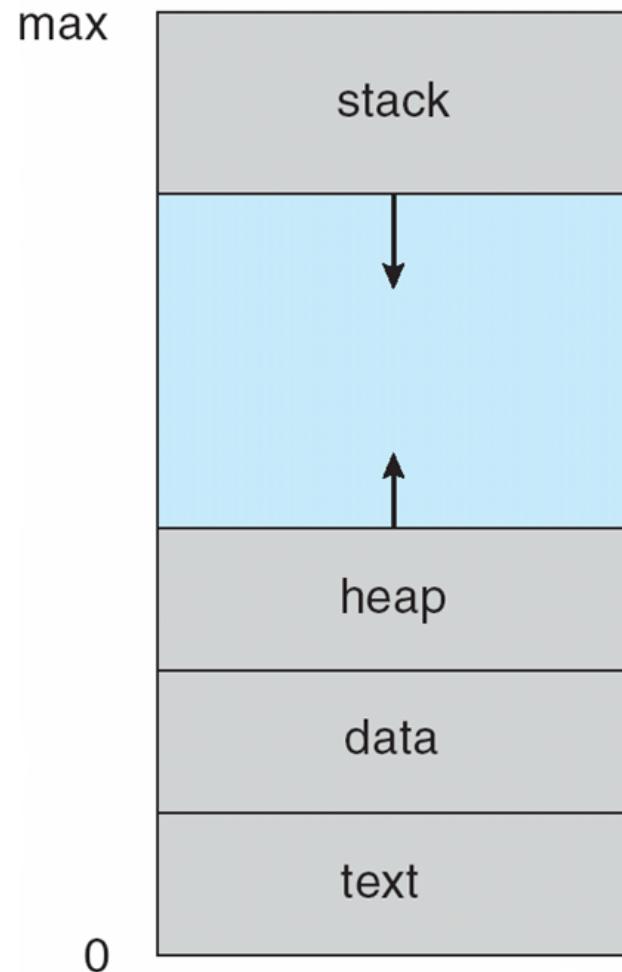
Objectives

- ▶ To introduce the notion of a process
- ▶ To describe the various features of processes, including scheduling, creation and termination, and communication
- ▶ To explore inter-process communication
- ▶ To describe communication in client-server systems

Basic Process Concept

- ▶ A program is a **passive** entity stored on disk, and a process is an **active** entity
 - A program becomes process when the executable file is loaded into memory
 - The execution of a program started via GUI mouse clicks, the command line entry of its name, etc.
 - One program can be executed as several processes
- ▶ An operating system can execute a variety of programs
 - In batch systems: jobs
 - In time-shared systems: user programs or tasks

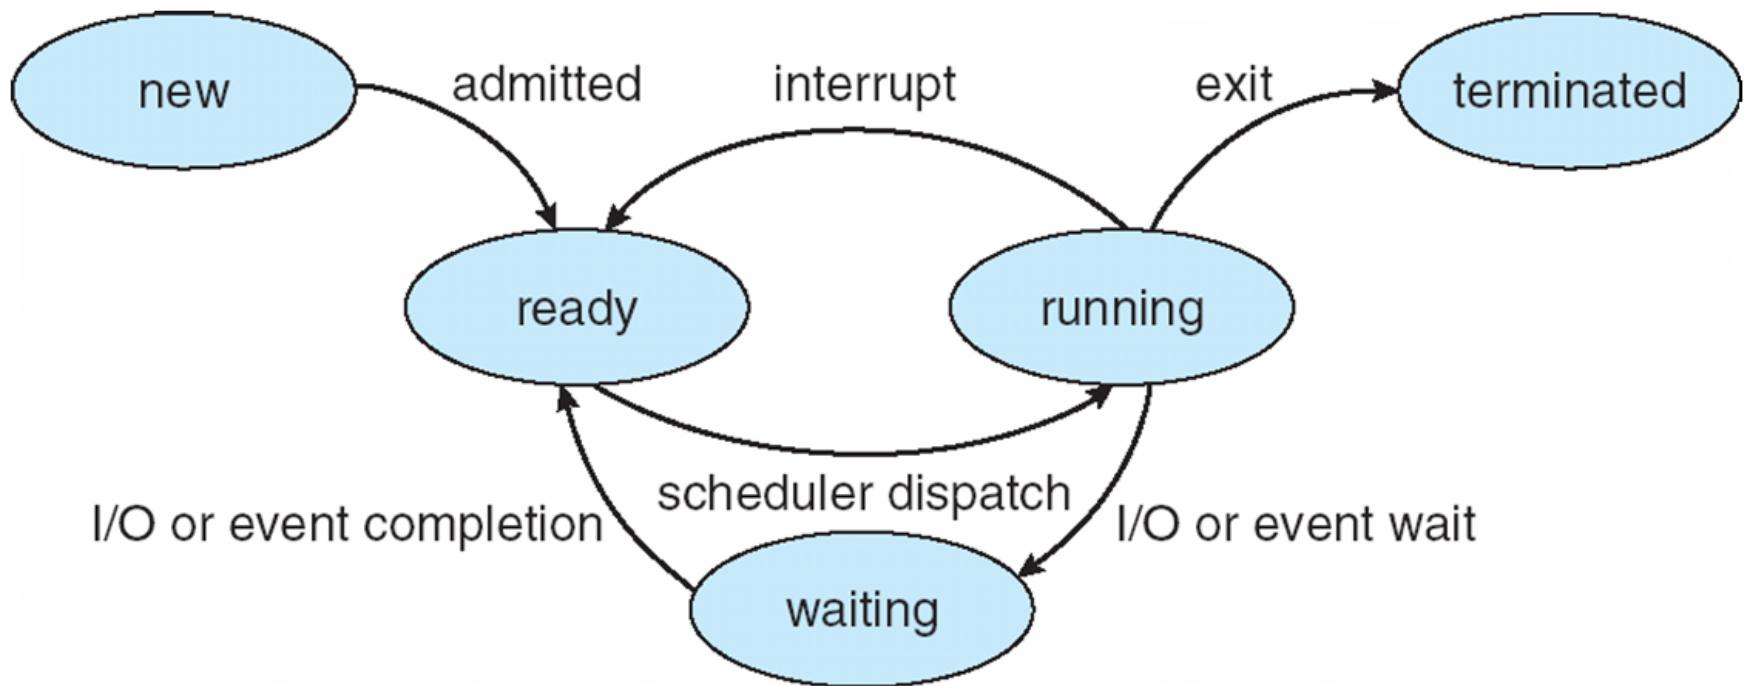
Process in Memory



Process States

- ▶ **New**: The process is being created
- ▶ **Running**: Instructions are being executed
- ▶ **Waiting**: The process is waiting for some event to occur
- ▶ **Ready**: The process is waiting to be assigned to a processor
- ▶ **Terminated**: The process has finished execution

Diagram of Process States

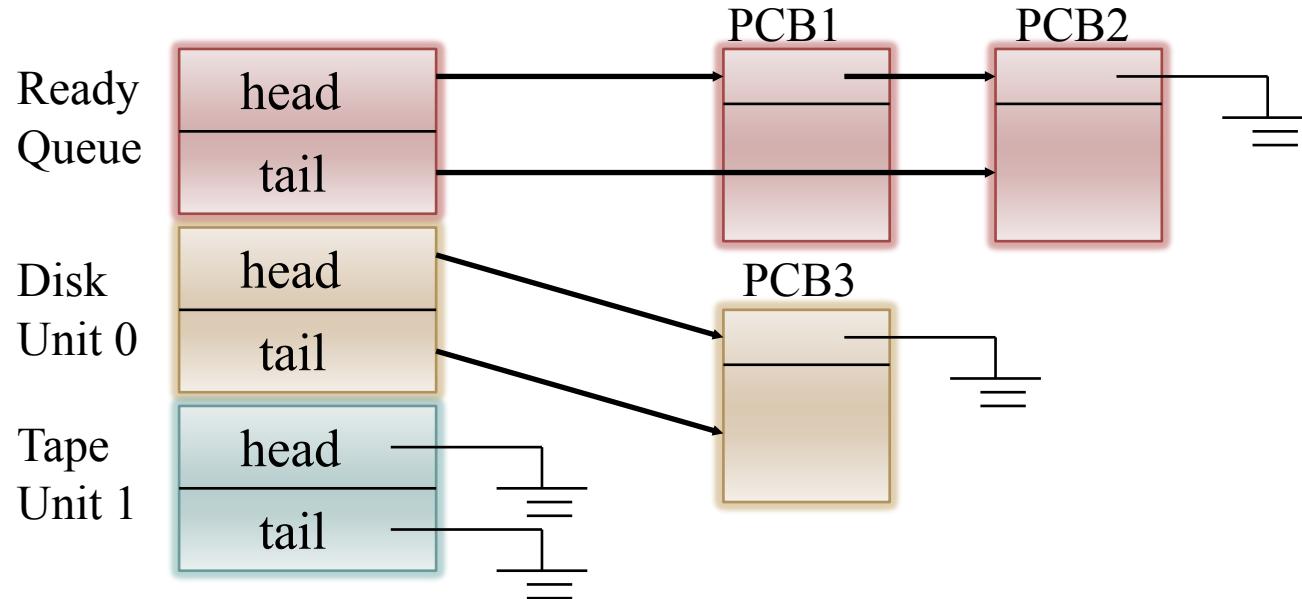


Process Control Block (PCB)

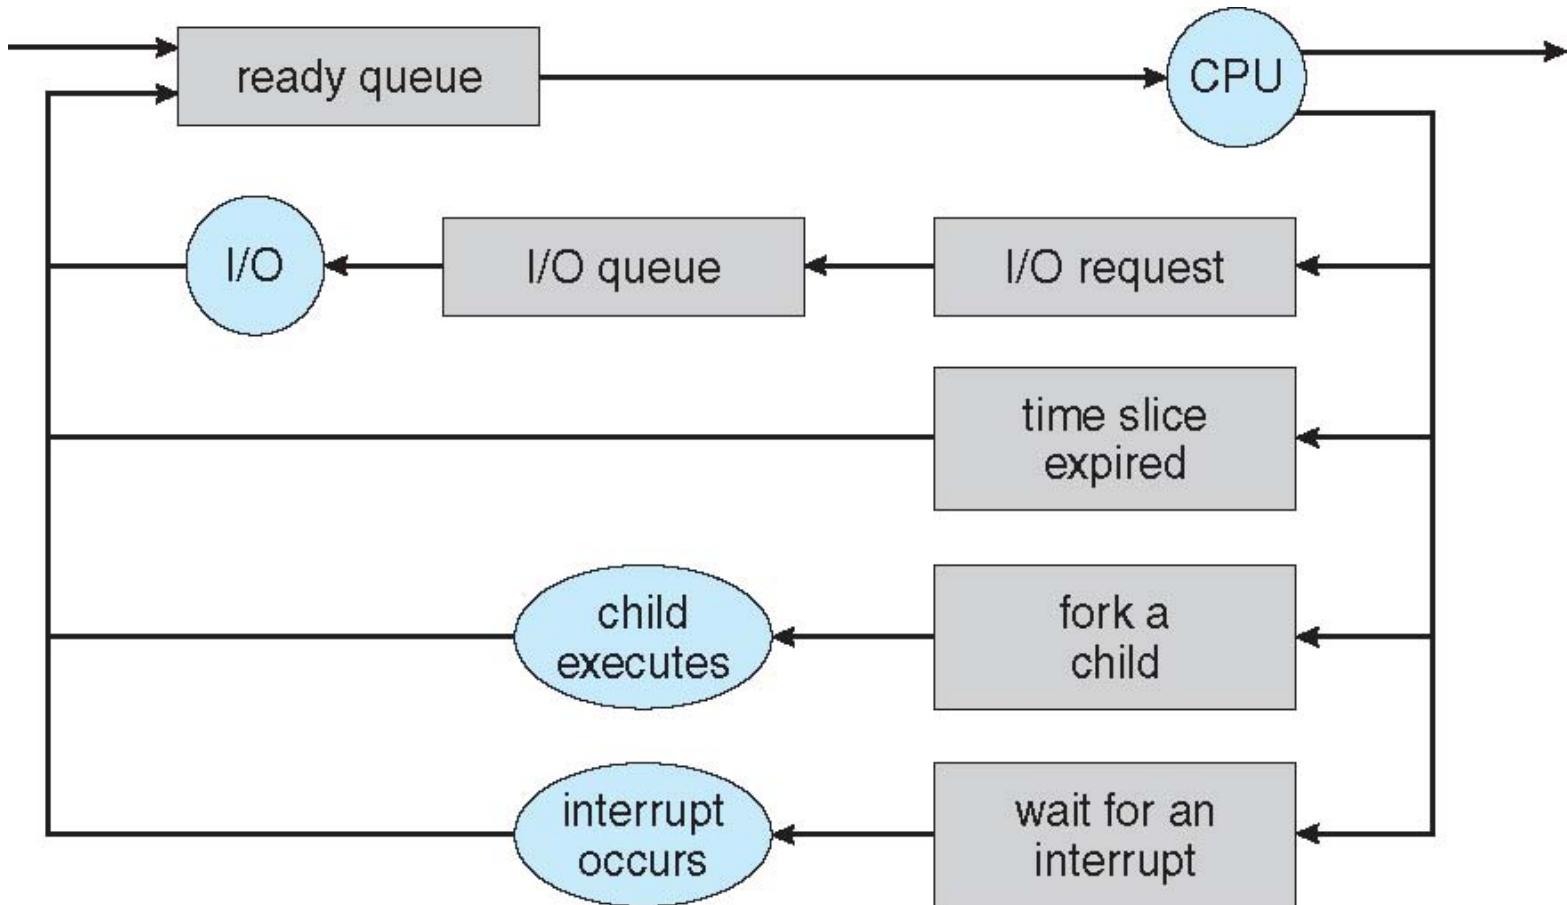
- ▶ PCB: The repository for any information that may vary from process to process
 - Process state— running, waiting, etc
 - Program counter— location of the currently executed instruction
 - CPU registers— contents of all process-centric registers
 - CPU scheduling information— priorities, scheduling queue pointers
 - Memory-management information— memory allocated to the process
 - Accounting information— CPU used, clock time elapsed since start, time limits
 - I/O status information— I/O devices allocated to process, list of opened files

Process Scheduling with PCB

- ▶ The goal of multiprogramming
 - Maximize CPU/resource utilization
- ▶ The goal of time sharing
 - Allow each user to interact with his/her program



Process Scheduling- A Queueing Diagram

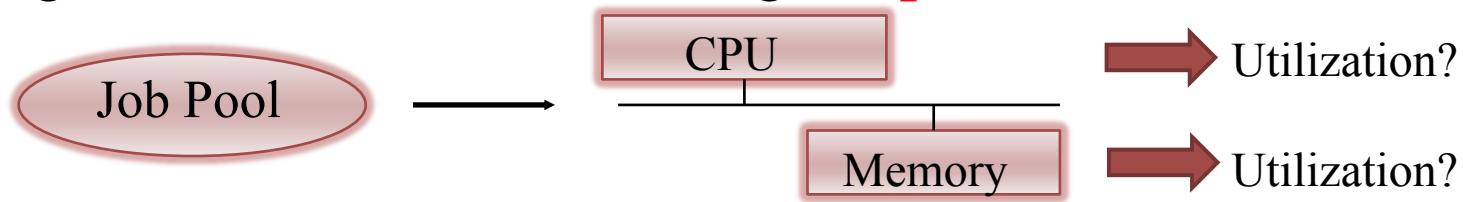


Processor Schedulers

- ▶ **Long-term scheduler (or job scheduler)**— selects which processes should be brought into the ready queue
- ▶ **Short-term scheduler (or CPU scheduler)**— selects which process should be executed next and allocates CPU
- ▶ **Medium-term scheduler** can be added as **swapper**

Long-Term Scheduler

- ▶ Processes can be described as either:
 - **I/O-bound process** – spends more time doing I/O than computations, many short CPU bursts
 - **CPU-bound process** – spends more time doing computations; few very long CPU bursts
- ▶ Long-term scheduler strives for good **process mix**



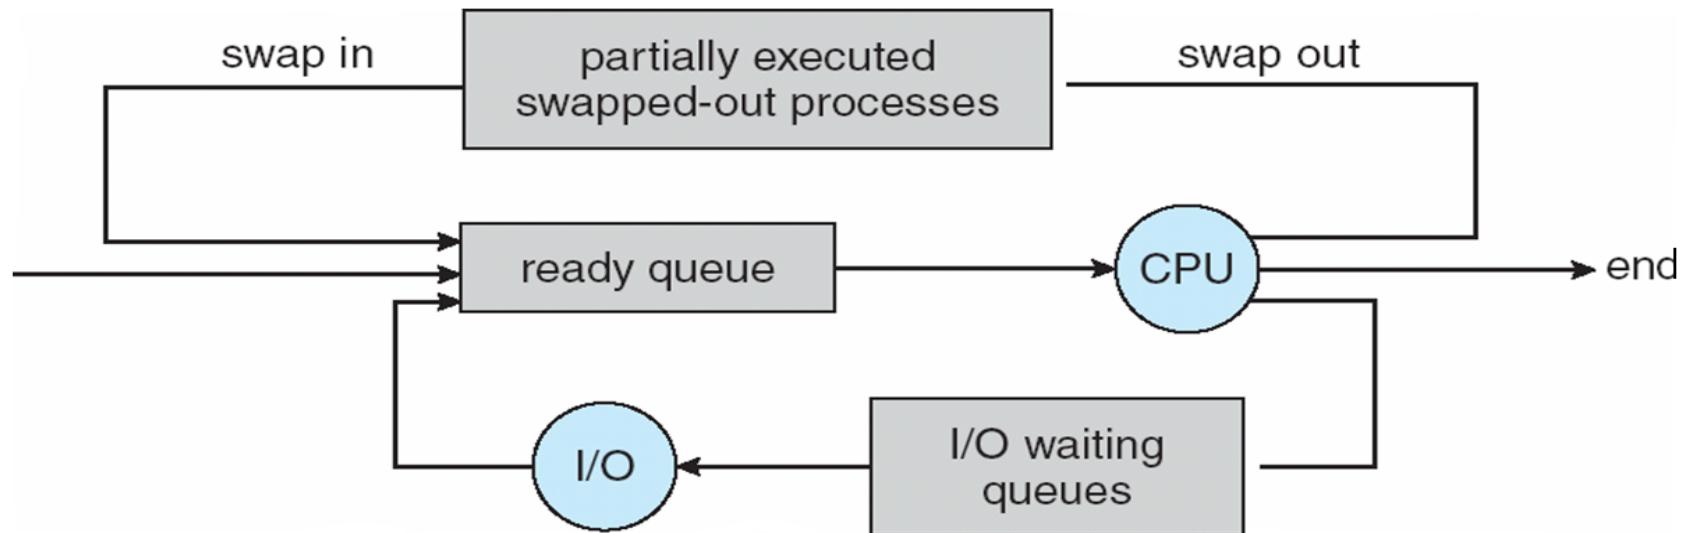
- ▶ Remarks :
 - Control the degree of multiprogramming
 - Can take more time in selecting processes because of a longer interval between executions
 - May not exist physically

Short-Term Scheduler

- ▶ Goal: To efficiently allocate the CPU to one of the ready processes according to some criteria
- ▶ Short-term scheduler is invoked very frequently (milliseconds) → must be fast
- ▶ In Linux, after version 2.6.23, the scheduler is the Completely Fair Scheduler (CFS)

Medium-Term Scheduler

- ▶ Goal: Remove process from memory, store on disk, bring back in from disk to continue execution: it is also called “swapping”



Process Scheduling- Context Switches

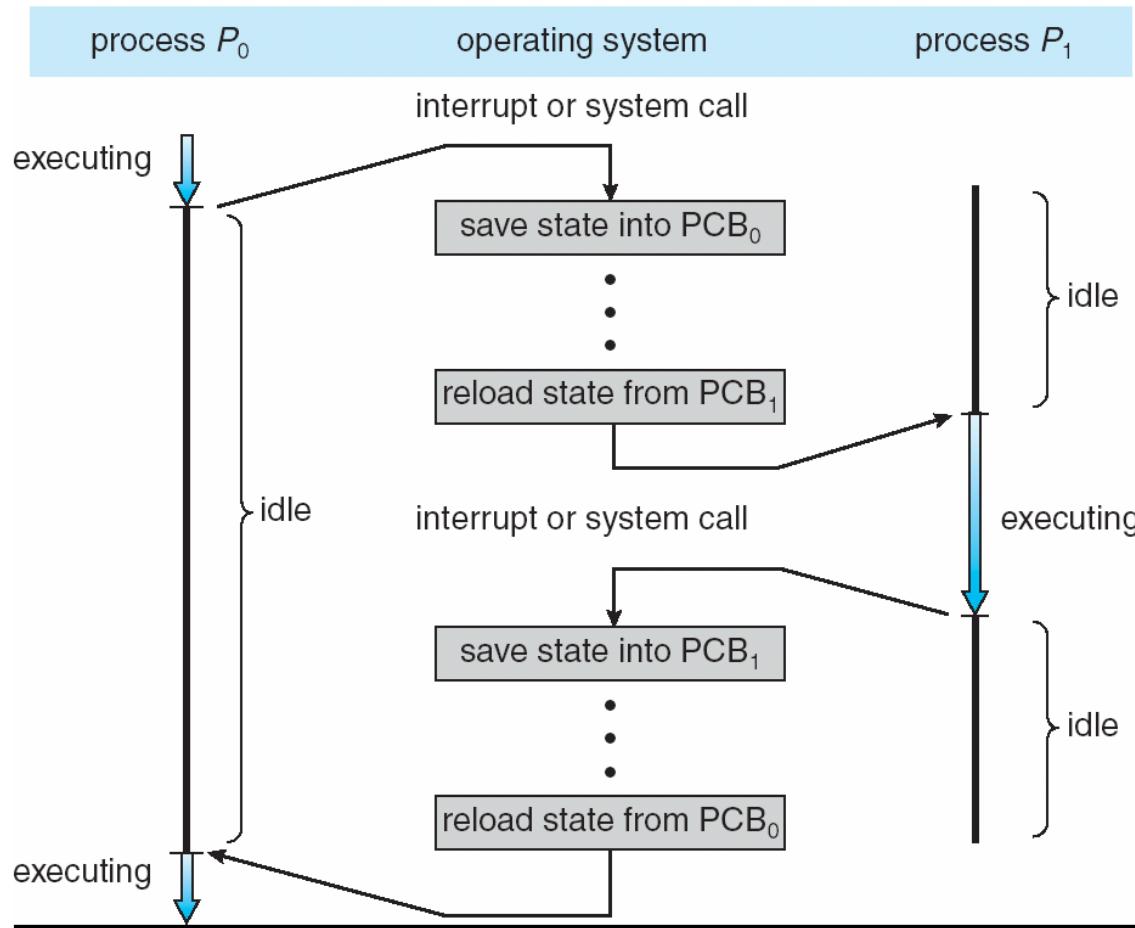
► Context Switch: Pure Overheads

- Save the state of the old process and load the state of the newly scheduled process.
 - The context of a process is usually reflected in PCB

► Issues:

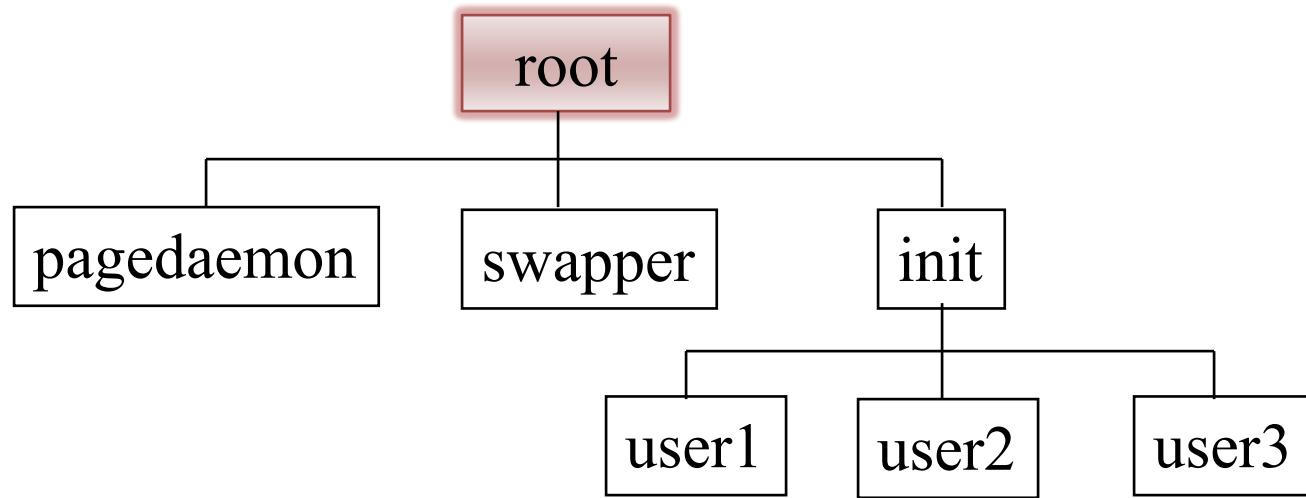
- The cost depends on the hardware support
 - e.g. processors with multiple register sets or computers with advanced memory management
- Threads, i.e., light-weight process (LWP), are introduced to break this bottleneck

CPU Switch from Process to Process



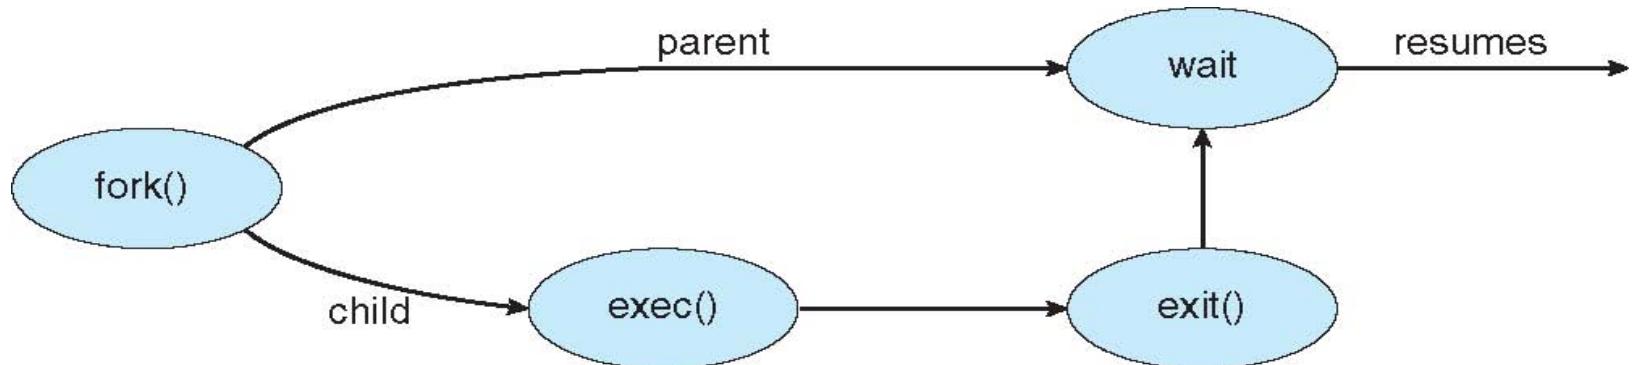
Parent and Child Processes

- ▶ Parent processes create child processes, which in turn create other processes, forming a tree of processes
- ▶ Generally, process identified and managed via a process identifier (PID)



Process Creation

- ▶ Address Space
 - Child duplicate of parent
 - Child has a program loaded into it
- ▶ UNIX Examples
 - **fork()** system call creates new process
 - **exec()** system call used after a **fork()** to replace the process' memory space with a new program



Process Termination

- ▶ Process executes last statement and asks the operating system to delete it: **exit()**
 - Wait the output data from child to parent: **wait()**
- ▶ Parent may terminate the execution of child processes:
abort()
 - ➔ Child has exceeded allocated resources
 - ➔ Task assigned to child is no longer required
 - Receive the return value form child
 - Some operating systems do not allow child to continue if its parent terminates
 - All children should be terminated - **cascading termination**

C Program Forking a Process

```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>

int main()
{
    pid_t pid;

    /* fork a child process */
    pid = fork();

    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;
    }
    else if (pid == 0) { /* child process */
        execlp("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child to complete */
        wait(NULL);
        printf("Child Complete");
    }

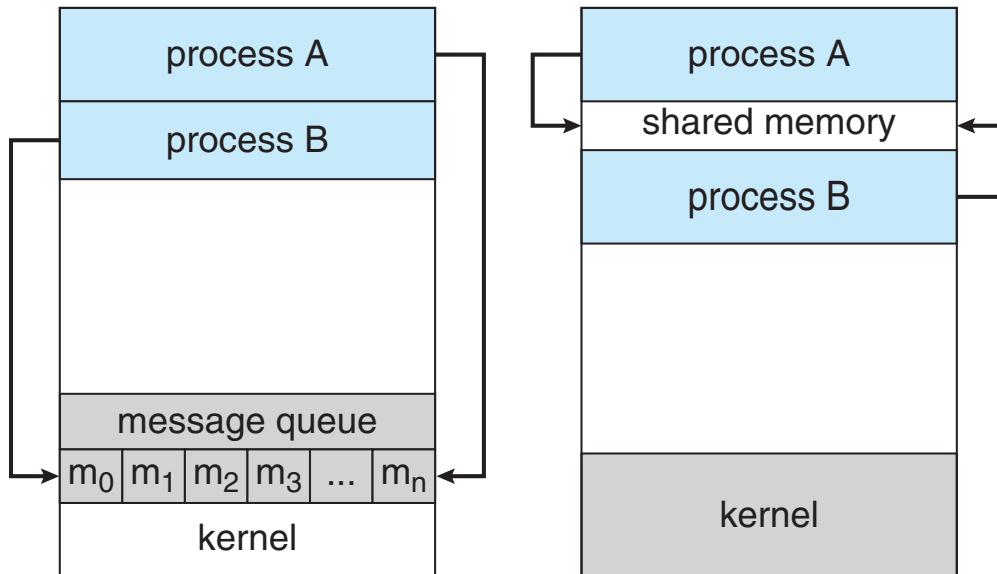
    return 0;
}
```

Inter-Process Communication

- ▶ Processes within a system may be independent or cooperating
- ▶ Cooperating process can affect or be affected by other processes, including sharing data
- ▶ Reasons for cooperating processes:
 - Information sharing
 - Computation speedup
 - Modularity
 - Convenience
- ▶ Cooperating processes need inter-process communication (IPC)

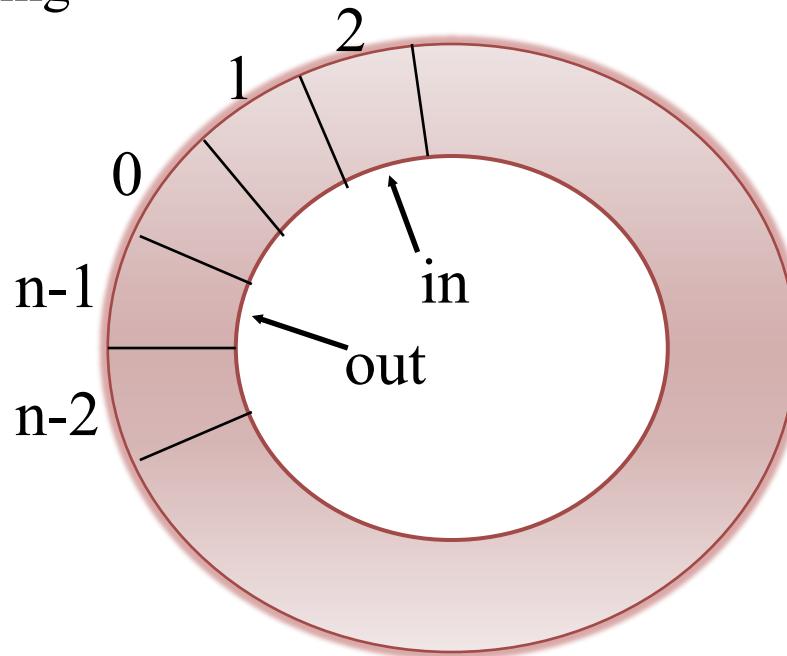
Two Models of IPC

- ▶ Shared Memory
 - Max Speed & Communication Convenience
- ▶ Message Passing
 - No Access Conflict & Easy Implementation



Shared Memory IPC

- ▶ A Consumer-Producer Example:
 - Bounded buffer or unbounded buffer
 - Supported by inter-process communication (IPC) or by hand coding



$\text{buffer}[0 \dots n-1]$

Initially, $\text{in}=\text{out}=0$

Shared Memory- Consumer

```
while (true)
{
    while (in == out);
    /* do nothing and have to wait */
    next_consumed = buffer[out];
    out = (out + 1) % BUFFER_SIZE;
    ... /* use the consumed item */
}
```

Shared Memory- Producer

```
while (true)
{
    ... /* produce a new item */
    while (((in + 1) % BUFFER SIZE) == out) ;
    /* do nothing */
    buffer[in] = next_produced;
    in = (in + 1) % BUFFER SIZE;
}
```

Message Passing IPC

- ▶ Logical Implementation of Message Passing
 - Fixed/variable message size
 - Symmetric/asymmetric communication
 - Direct/indirect communication
 - Synchronous/asynchronous communication
 - Automatic/explicit buffering
 - Send by copy or reference

Direct Message Passing

- ▶ Processes must name each other explicitly:
 - **send** ($P, message$) – send a message to process P
 - **receive** ($Q, message$) – receive a message from process Q
- ▶ Properties of the communication link
 - Links are established automatically
 - A link is associated with exactly one pair of communicating processes
 - Between each pair there exists exactly one link
 - The link may be unidirectional, but is usually bi-directional

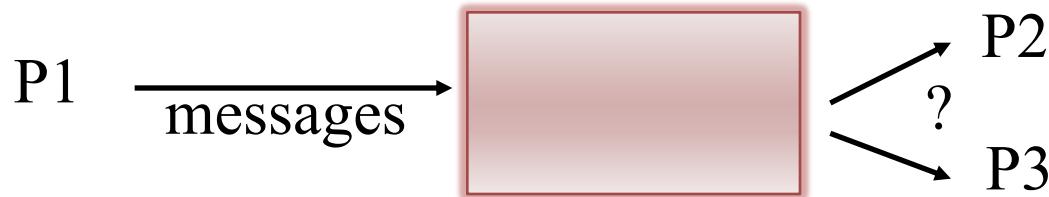
Indirect Communication

- ▶ Messages are directed and received from mailboxes (also referred to as ports)
 - **send**($A, message$) – send a message to mailbox A
 - **receive**($A, message$) – receive a message from mailbox A
- ▶ Properties of the communication link
 - Links are established only if processes share a common mailbox
 - A link may be associated with many processes
 - Each pair of processes may share several communication links
 - Links may be unidirectional or bi-directional

Issues of Indirect Communication

► Mailbox sharing

- P_1 , P_2 , and P_3 share mailbox A
- P_1 , sends; P_2 and P_3 receive
- Who gets the message?



► Solutions

- Allow a link to be associated with at most two processes
- Allow only one process at a time to execute a receive operation
- Allow the system to select arbitrarily the receiver

IPC Synchronization

- ▶ Synchronous Message Passing IPC
 - Blocking send has the sender block until the message is received
 - Blocking receive has the receiver block until a message is available
- ▶ Asynchronous Message Passing IPC
 - Non-blocking send has the sender send the message and continue
 - Non-blocking receive has the receiver receive a valid message or null

IPC Buffering

- ▶ The capacity of a link: the number of messages could be held in the link
 - Zero capacity – 0 messages
 - Sender must wait for receiver
 - Bounded capacity – finite length of n messages
 - Sender must wait if link is full
 - Unbounded capacity – infinite length
 - Sender never waits
- ▶ The last two items are for asynchronous communication and may need **acknowledgement**

Examples of IPC Systems – POSIX

▶ POSIX Shared Memory

- Process first creates shared memory segment

```
shm_fd = shm_open(name, O_CREAT | O_RDWR,  
0666);
```

- Set the size of the object

```
ftruncate(shm_fd, 4096);
```

- Memory map the object

```
ptr = mmap(0, 4096, PROT_WRITE, MAP_SHARED,  
shm_fd, 0);
```

- Now the process could write to the shared memory

```
sprintf(ptr, "Writing to shared memory");
```

Examples of IPC Systems – Mach

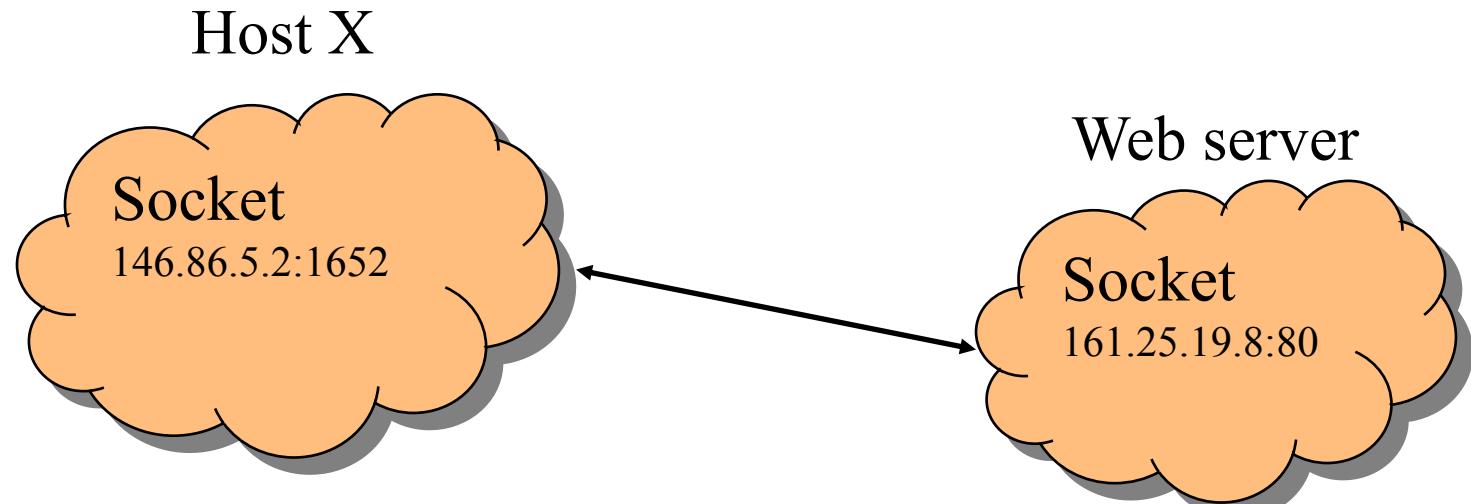
- ▶ Mach – A message-based OS from the Carnegie Mellon University
 - When a task is created, two special mailboxes, called ports, are also created.
 - The *Kernel* mailbox is used by the kernel to communicate with the tasks
 - The *Notify* mailbox is used by the kernel sends notification of event occurrences.

Three IPC System Calls on Mach

- ▶ **msg_send**
 - Options when mailbox is full:
 - Wait indefinitely
 - Return immediately
 - Wait at most for n ms
 - Temporarily cache a message: only one message to a full mailbox can be pending at any time for a sending tread
- ▶ **msg_receive**
 - Only one task can own & have a receiving privilege of a mailbox
 - Options when mailbox is empty:
 - Wait indefinitely
 - Return immediately
 - Wait at most for n ms
- ▶ **msg_rpc**
 - Remote Procedure Calls

Communication in Client–Server Systems

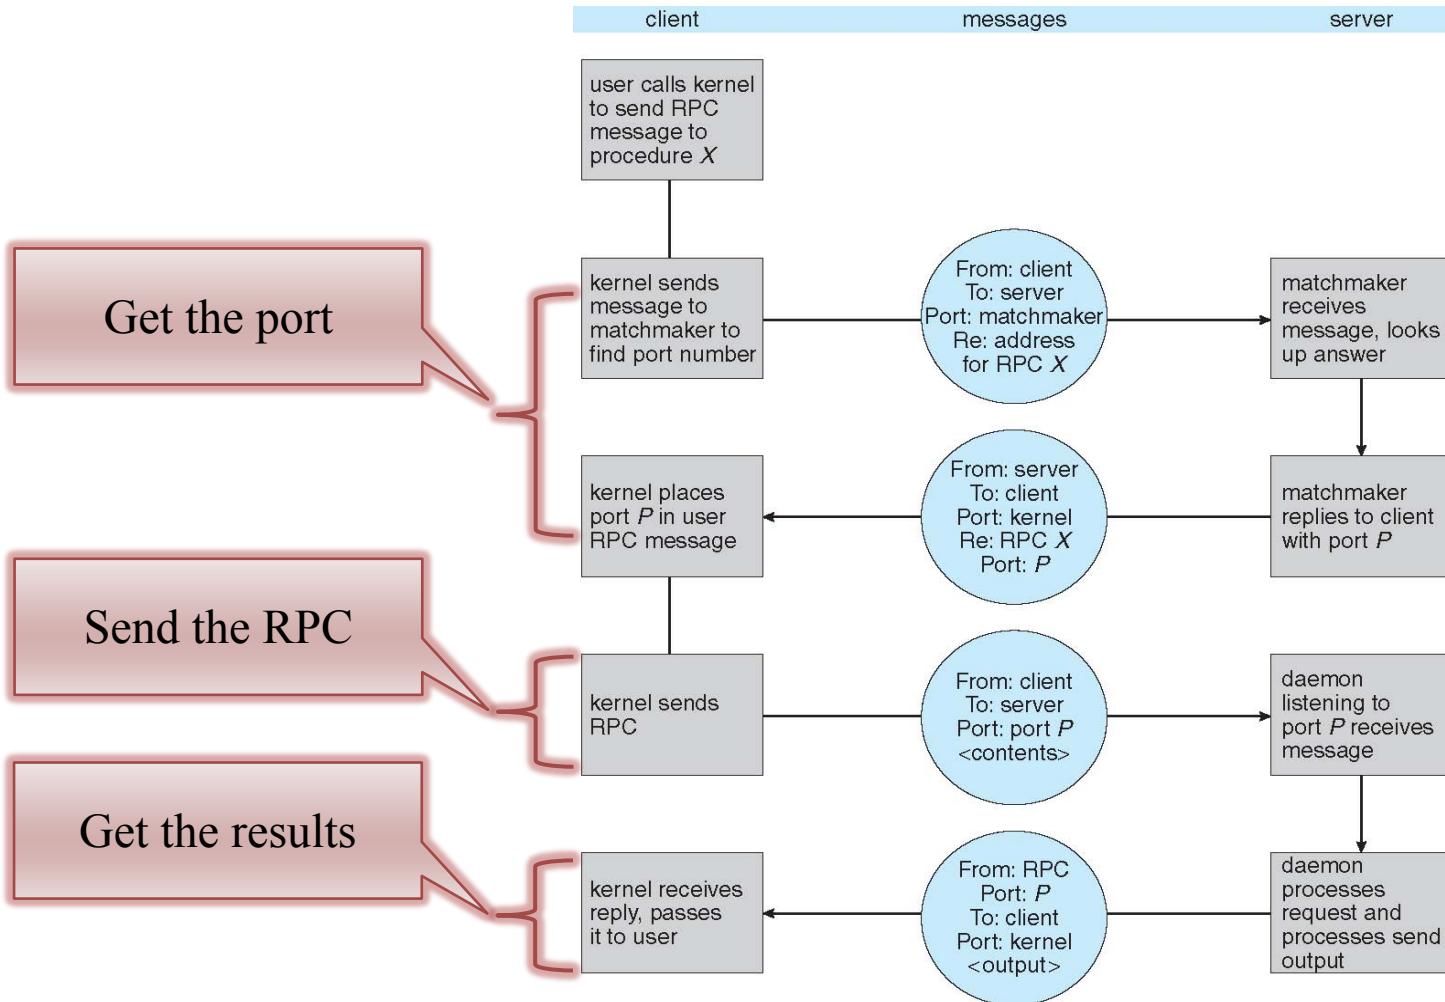
- ▶ Socket
 - An endpoint for communication identified by an IP address concatenated with a port number
 - A client-server architecture
- ▶ /etc/services: 23-telnet, 21-ftp, 80-web server, etc.



Remote Procedure Calls

- ▶ A way to abstract the procedure-call mechanism for use between systems with network connection
- ▶ Stubs at the client site
 - One for each RPC
 - Locate the proper port and marshal parameters
- ▶ Stubs at the server site
 - Receive the message
 - Invoke the procedure and return the results
- ▶ Data representation handled via the External Data Representation (XDL) format to account for different architectures
 - Big-endian and little-endian

Execution of RPC

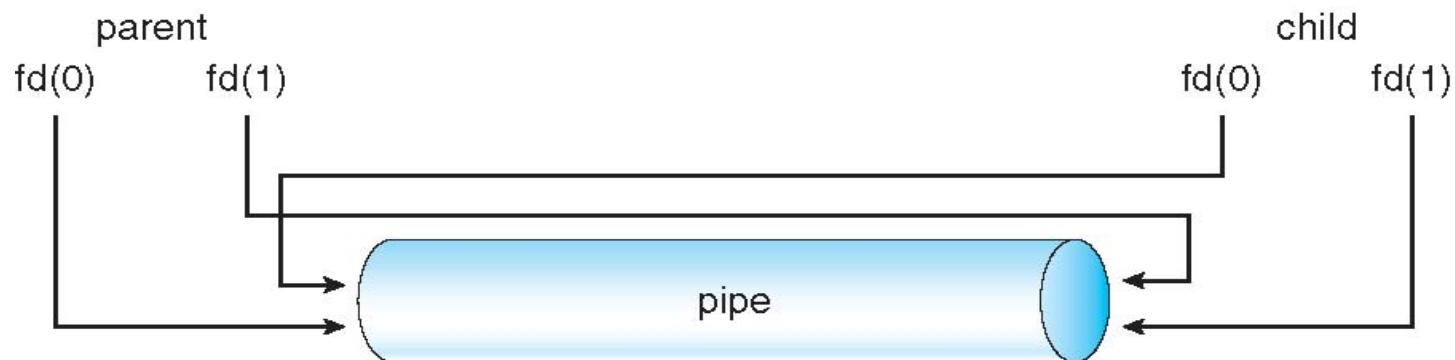


Pipes

- ▶ Acts as a conduit allowing two processes to communicate
- ▶ Issues
 - Is communication unidirectional or bidirectional?
 - In the case of two-way communication, is it half-duplex or full-duplex?
 - Must there exist a relationship (i.e. parent-child) between the communicating processes?
 - Can the pipes be used over a network?

Ordinary Pipes

- ▶ Ordinary Pipes allow communication in the standard producer-consumer style
- ▶ Producer writes to the **write-end** of the pipe
- ▶ Consumer reads from the **read-end** of the pipe
- ▶ Ordinary pipes are therefore unidirectional
- ▶ Require parent-child relationship between communicating processes



Named Pipes

- ▶ Named Pipes are more powerful than ordinary pipes
- ▶ Communication is bidirectional
- ▶ No parent-child relationship is necessary between the communicating processes
- ▶ Several processes can use the named pipe for communication
- ▶ Provided on both UNIX and Windows systems